PhotoMind: RL-Powered Multimodal Personal Photo Knowledge Retrieval

A Multi-Agent System with Contextual Bandits and
DQN Confidence Calibration

Course: Prompt Engineering — Building Agentic Systems
Assignment: Take-Home Final — Reinforcement Learning for Agentic AI
Platform: CrewAI (Python) · PyTorch · React · FastAPI

Raghunandan Gunji
Northeastern University
gunji.r@northeastern.edu

April 22, 2026

Abstract

PhotoMind is a multimodal personal photo knowledge retrieval system that transforms a user's photo library into a queryable knowledge base. The system uses four specialized CrewAI agents—Controller, Photo Analyst, Knowledge Retriever, and Insight Synthesizer—orchestrated via hierarchical delegation to answer natural-language queries over personal photos (receipts, food, documents, screenshots). We extend the base system with two reinforcement learning components: (1) a contextual bandit (Thompson Sampling, UCB1, ε-greedy) that routes queries across four search strategies (factual, semantic, behavioral, embedding), and (2) a Deep Q-Network (DQN) confidence calibrator with a non-terminal *requery* action that enables multi-step evidence gathering. Both components are trained offline via a deterministic simulator at $0 API cost. A 7-configuration ablation study across 5 random seeds with 95% confidence intervals demonstrates that the DQN eliminates silent failures (confident-but-wrong answers) from 2.4% to 0.0% on training queries, with the Recommended (Rule+DQN) configuration achieving 0.4% silent failure rate while preserving baseline routing accuracy. The system includes a React frontend with FastAPI backend, dense vector embedding search via sentence-transformers (all-MiniLM-L6-v2), and comprehensive evaluation infrastructure with 56 queries, statistical significance testing, and held-out generalization validation.

Contents

1 Introduction **5**
- 1.1 Core Capabilities . . . 5
- 1.2 Real-World Use Cases . . . 5
- 1.3 Project Scope and Contributions . . . 5

2 System Architecture **5**
- 2.1 High-Level Architecture . . . 5
- 2.2 Data Flow . . . 6
- 2.3 Architecture with RL Components . . . 10
- 2.4 Memory Architecture . . . 12

3 Agent Design **12**
- 3.1 Controller Agent (Query Crew Manager) . . . 12
- 3.2 Photo Analyst (Ingestion Crew) . . . 12
- 3.3 Knowledge Retriever . . . 12
- 3.4 Insight Synthesizer . . . 12
- 3.5 Agent Collaboration Pattern . . . 12

4 Tool Integration **13**
- 4.1 Built-in Tools . . . 13
- 4.2 Custom Vision Tool (`PhotoVisionTool`) . . . 13
- 4.3 Custom Knowledge Base Tool (`PhotoKnowledgeBaseTool`) . . . 13
  - 4.3.1 Input/Output Schema . . . 13

5 Search Strategies **13**
- 5.1 Factual Search . . . 14
- 5.2 Semantic Search . . . 14
- 5.3 Behavioral Search . . . 14
- 5.4 Embedding Search (New) . . . 14
- 5.5 Confidence Grading . . . 15

6 Orchestration Design **15**
- 6.1 Ingestion Crew (Sequential) . . . 15
- 6.2 Query Crew (Hierarchical with Planning) . . . 15
- 6.3 Feedback Loop . . . 15
- 6.4 Communication Protocols . . . 15

7 Reinforcement Learning Extension **16**
- 7.1 Motivation . . . 16
- 7.2 RL Approach 1: Contextual Bandits . . . 16
  - 7.2.1 Problem Formulation . . . 16
  - 7.2.2 Context Clustering . . . 16
  - 7.2.3 Feature Extractor . . . 16
  - 7.2.4 Cluster Count Selection ($k = 4$) . . . 16
  - 7.2.5 Thompson Sampling . . . 17
  - 7.2.6 UCB1 . . . 17

7.2.7 Bandit Reward Function . . . 17
7.3 RL Approach 2: DQN Confidence Calibrator . . . 18
7.3.1 MDP Formulation . . . 18
7.3.2 Multi-Step Bellman Equation . . . 18
7.3.3 Q-Network Architecture . . . 18
7.3.4 TD Learning . . . 18
7.3.5 Reward Matrix . . . 19
7.3.6 Hyperparameters . . . 19
7.4 Offline Simulation Environment . . . 19

8 Base System Performance **20**
8.1 Evaluation Setup . . . 20
8.2 Results . . . 20
8.3 Per-Category Breakdown . . . 20

9 RL Experimental Results **21**
9.1 Training Convergence . . . 21
9.2 Primary Result: Recommended Configuration . . . 21
9.3 Ablation Study . . . 21
9.4 Statistical Significance . . . 22
9.5 Learning Dynamics Visualizations . . . 23
9.6 Held-Out Generalization . . . 24
9.7 Cross-Seed Variance Analysis . . . 25

10 Discussion **25**
10.1 When RL Helps . . . 25
10.2 Multi-Step Requery Dynamics . . . 25
10.3 Per-Category RL Impact . . . 25
10.4 Bandit × DQN Interaction Effects . . . 26
10.5 Routing Accuracy Regression Analysis . . . 26
10.6 Why Rule+DQN Outperforms Full RL . . . 26
10.7 Cost Analysis . . . 26

11 Scaling Analysis **27**

12 React Frontend and FastAPI Backend **27**
12.1 Architecture . . . 27
12.2 Dual Query Modes . . . 27
12.3 API Endpoints . . . 28

13 Challenges and Solutions **28**

14 Threats to Validity **28**
14.1 Internal Validity . . . 28
14.2 External Validity . . . 29
14.3 Construct Validity . . . 29
14.4 Statistical Conclusion Validity . . . 29

15 Ethical Considerations **29**

16 Future Work **30**
16.1 Immediate Extensions 30
16.2 Production Deployment Roadmap 30
16.3 Research Directions 30

17 Reproducibility **30**

18 Conclusion **31**

1 Introduction

Smartphone users accumulate thousands of photos—receipts, bills, food, screenshots, documents—but cannot search them by meaning or extract facts using natural language. PhotoMind addresses this gap by building a structured knowledge base from photo analysis and providing confidence-graded answers with source attribution.

1.1 Core Capabilities

- Analyze photos using GPT-4o Vision to extract text (OCR), entities, and semantic descriptions
- Build a persistent structured knowledge base from the analysis results
- Answer natural-language queries with four distinct retrieval strategies
- Return confidence-graded answers (A–F) with source photo attribution
- Gracefully decline queries the knowledge base cannot answer

1.2 Real-World Use Cases

- *"How much did I spend at ALDI?"* — finds all 5 ALDI receipts, aggregates total with source photos
- *"Show me photos of pizza"* — finds food photos matching the description
- *"What type of food do I photograph most?"* — analyzes patterns across all 25 photos
- *"What was my electric bill?"* — correctly declines (no such photo in the library)

1.3 Project Scope and Contributions

This project makes four principal contributions:

1. **Multi-component ablation methodology** that isolates individual and interaction effects of modular RL components, revealing that upstream stochasticity compounds into downstream variance.
2. **Zero-cost offline training** via deterministic simulation—a transferable pattern for any agentic system with expensive API calls.
3. **Asymmetric reward design** that prioritizes safety over coverage in personal-data domains.
4. **Requery as a non-terminal MDP action**, transforming single-step confidence classification into multi-step active information gathering with structurally relevant discounting.

2 System Architecture

2.1 High-Level Architecture

PhotoMind consists of two CrewAI crews: a **sequential ingestion crew** that analyzes photos and builds the knowledge base, and a **hierarchical query crew** that answers natural-language questions. Figure 1 shows the overall system architecture.

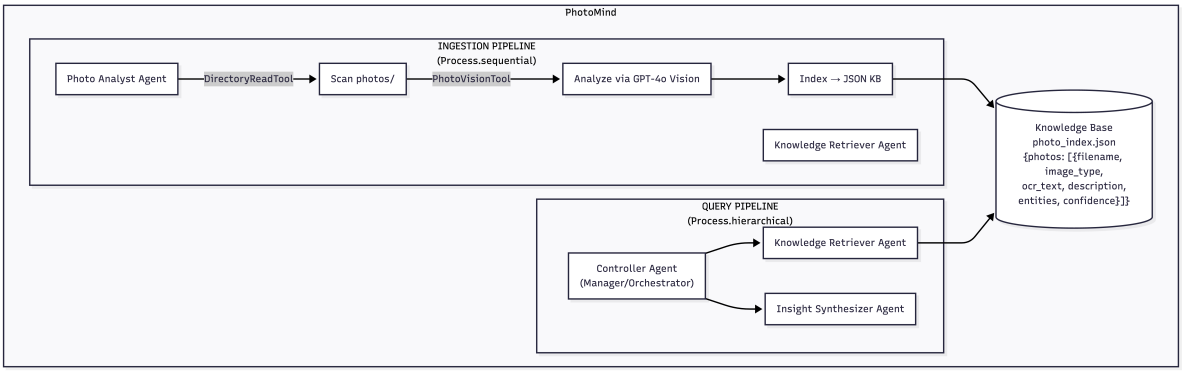

Figure 1: High-level system architecture showing ingestion and query pipelines.

2.2 Data Flow

Ingestion Pipeline. The sequential crew processes photos in three stages: (1) directory scan, (2) GPT-4o Vision analysis extracting OCR text, entities, descriptions, and confidence scores, (3) knowledge base indexing into a structured JSON file. The pipeline is idempotent—re-running on a partially indexed corpus only processes new photos.

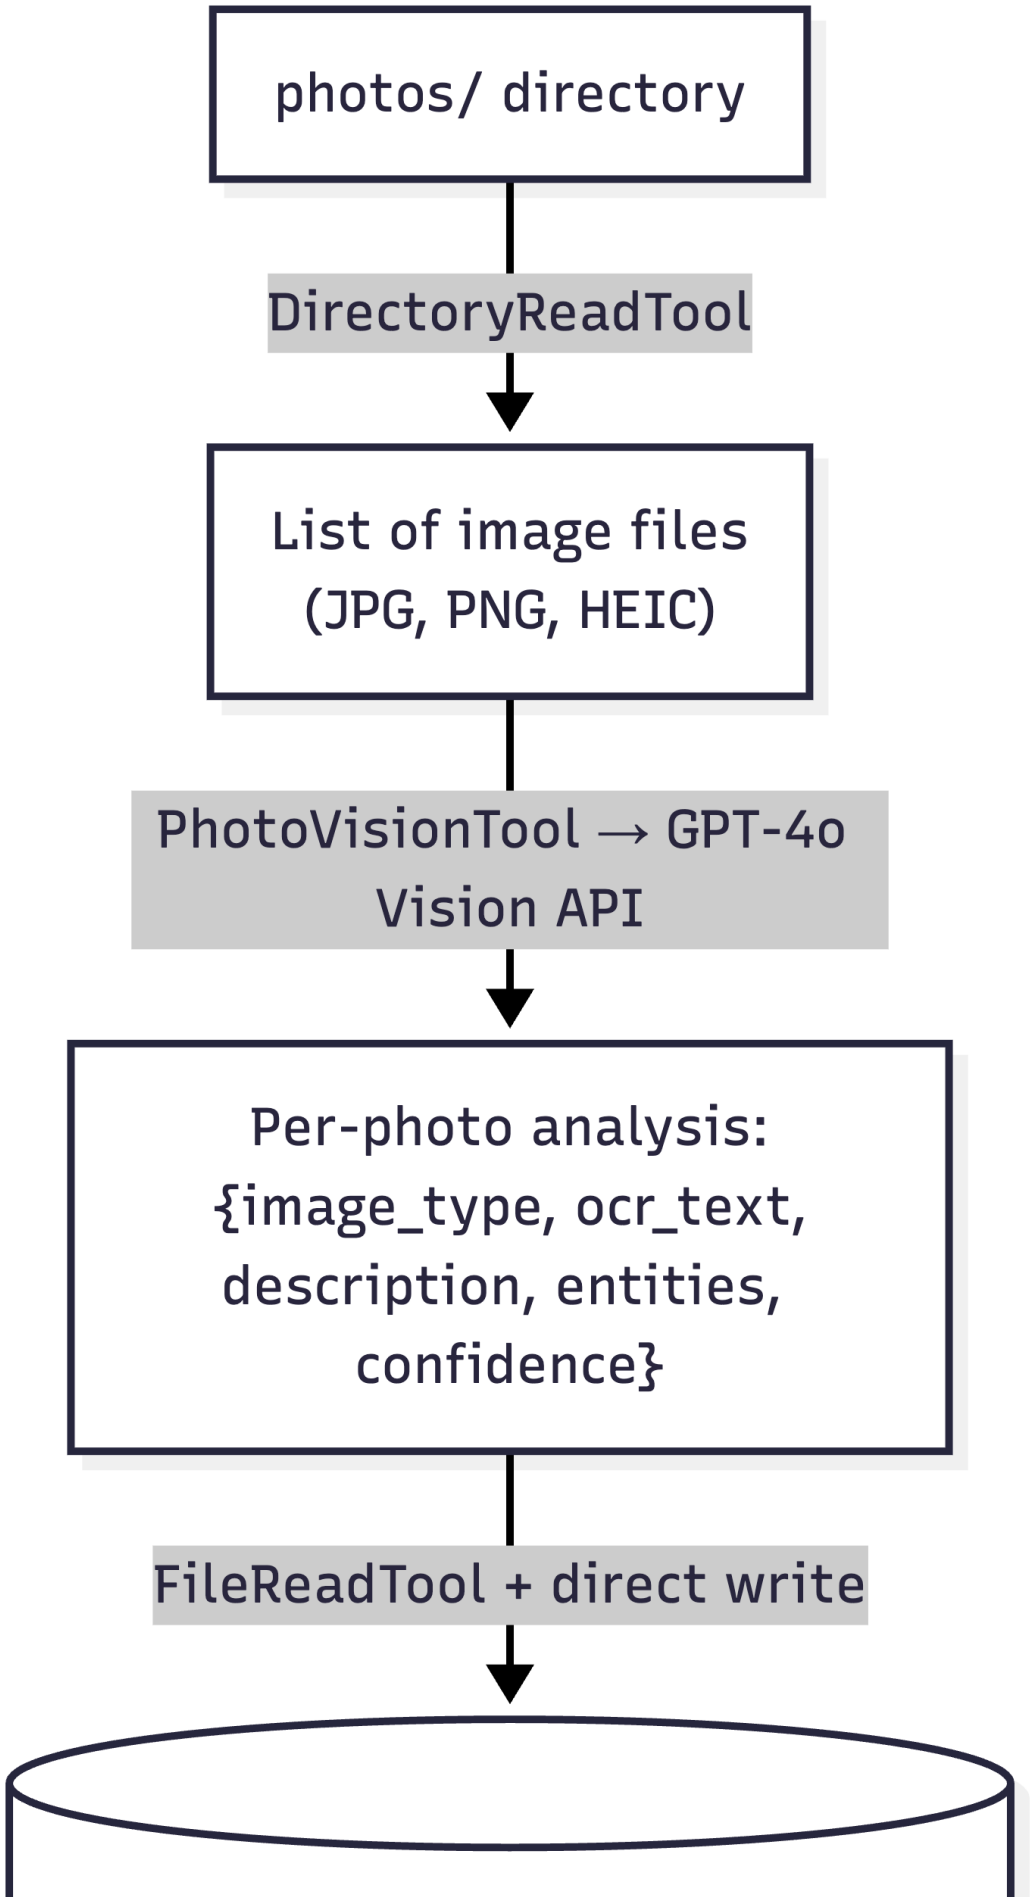
photos/ directory
DirectoryReadTool
List of image files (JPG, PNG, HEIC)
PhotoVisionTool → GPT-4o Vision API
Per-photo analysis: {image_type, ocr_text, description, entities, confidence}
FileReadTool + direct write

Query Pipeline. The hierarchical crew uses a Controller agent as the manager, delegating to the Knowledge Retriever (which uses the RL-enhanced PhotoKnowledgeBaseTool) and then the Insight Synthesizer for answer generation.

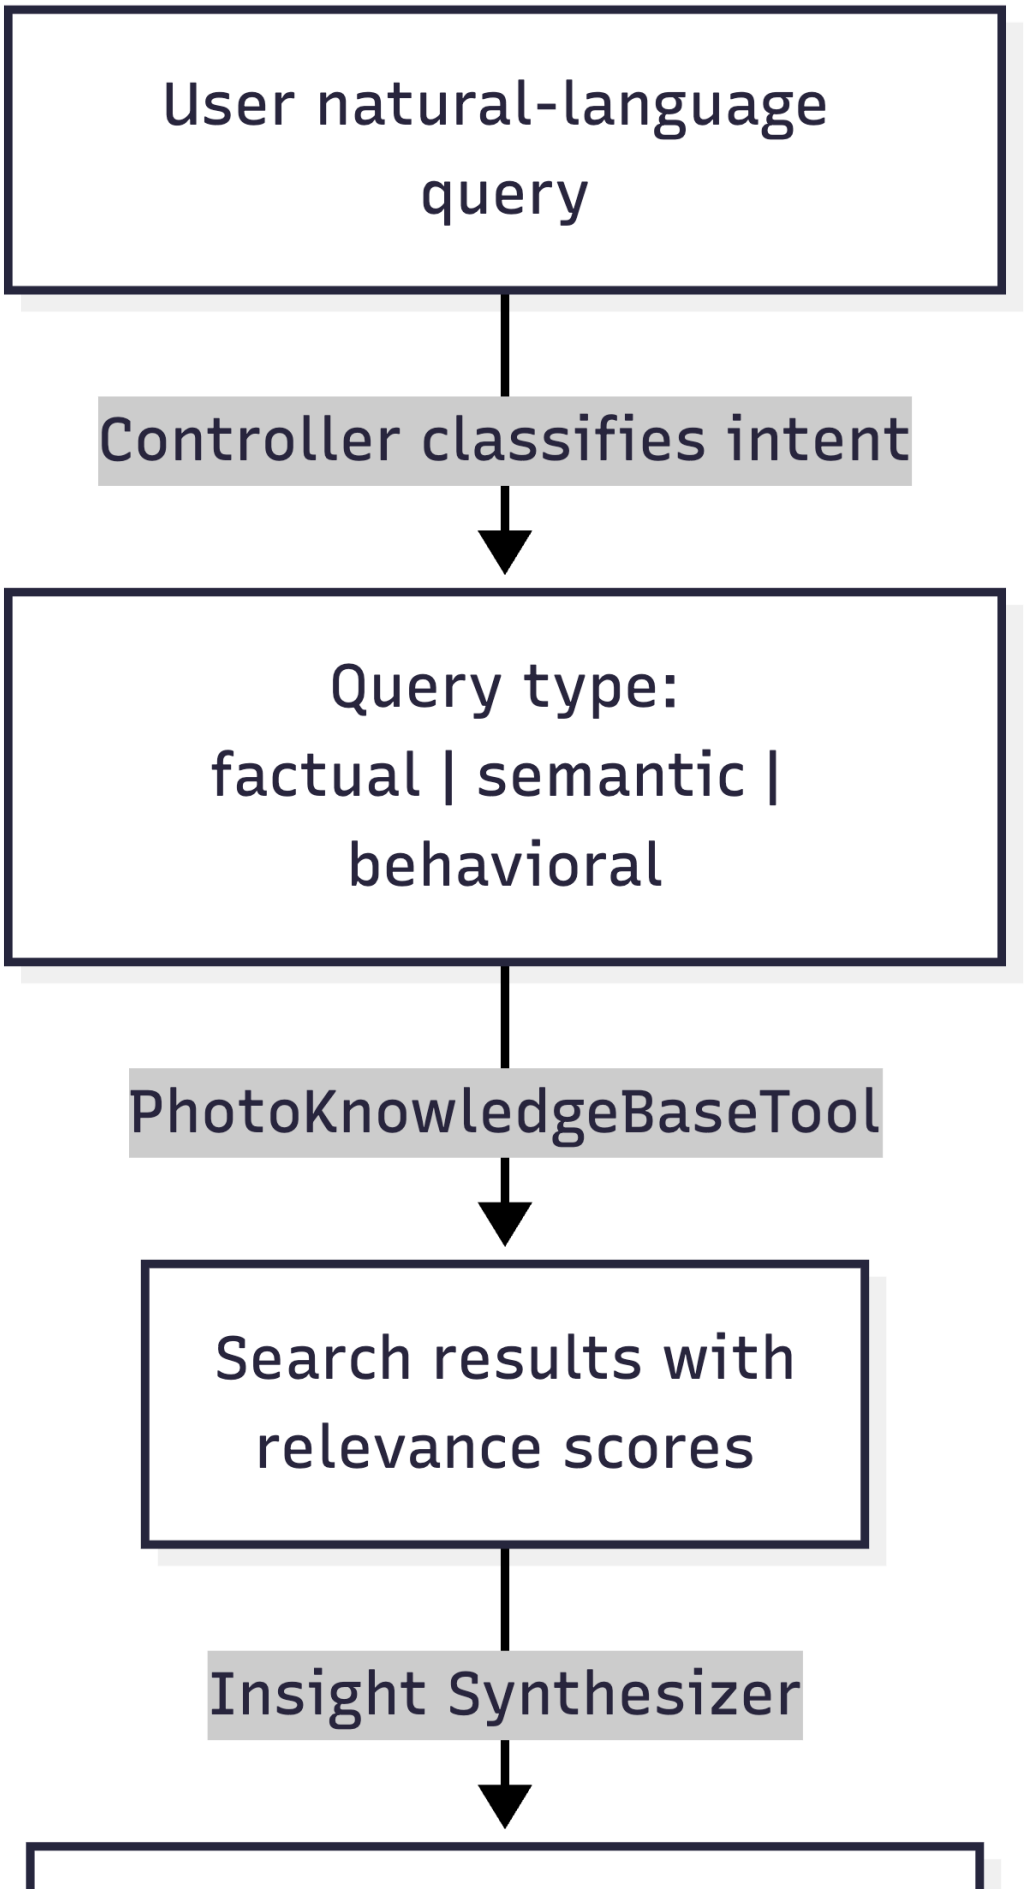
User natural-language query
Controller classifies intent
Query type: factual | semantic | behavioral
PhotoKnowledgeBaseTool
Search results with relevance scores
Insight Synthesizer

2.3 Architecture with RL Components

The RL extension adds two learned components to the query pipeline: a contextual bandit for strategy routing and a DQN for confidence calibration. Figure 4 shows the integrated architecture.

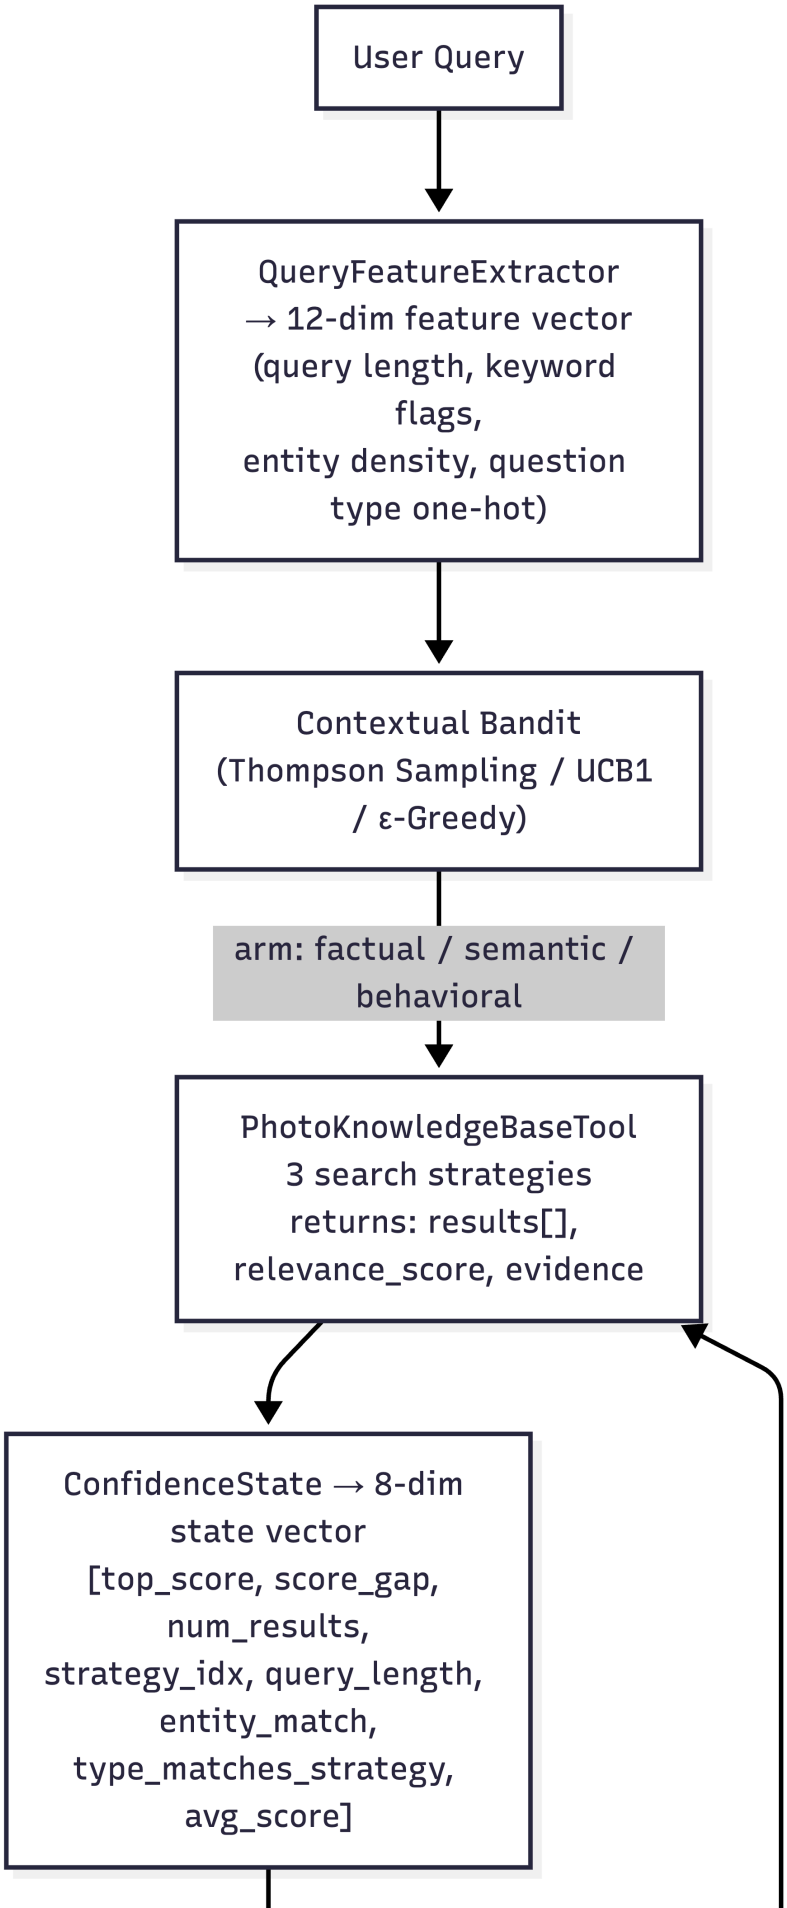
User Query
QueryFeatureExtractor
→ 12-dim feature vector
(query length, keyword
flags,
entity density, question
type one-hot)
Contextual Bandit
(Thompson Sampling / UCB1
/ ε-Greedy)
arm: factual / semantic /
behavioral
PhotoKnowledgeBaseTool
3 search strategies
returns: results[],
relevance_score, evidence
ConfidenceState → 8-dim
state vector
[top_score, score_gap,
num_results,
strategy_idx, query_length,
entity_match,
type_matches_strategy,
avg_score]

2.4 Memory Architecture

Both pipelines use CrewAI's built-in memory with a local sentence-transformer model (`all-MiniLM-L6-v2`) for vector embeddings, storing agent interactions in a local ChromaDB vector store for contextual awareness across agent steps without requiring external embedding API access.

3 Agent Design

PhotoMind uses four specialized agents across two crews.

3.1 Controller Agent (Query Crew Manager)

The Controller acts as the chief librarian: it classifies query intent (factual, semantic, behavioral) and delegates to the appropriate specialist agent. It operates as the manager in CrewAI's hierarchical process with `allow_delegation=True`.

Backstory: "You are the chief librarian of a personal photo knowledge base. You classify first, then delegate. If you are unsure, you say so."

3.2 Photo Analyst (Ingestion Crew)

Scans the `photos/` directory and passes each image to GPT-4o Vision for structured extraction: image classification, OCR text, semantic description, and named entities (amounts, dates, vendors, locations). Reports partial extraction with lower confidence rather than guessing on unclear images.

3.3 Knowledge Retriever

In ingestion: writes the knowledge base JSON. In queries: executes the search using the RL-enhanced `PhotoKnowledgeBaseTool`. The query task restricts this agent to only the primary search tool via task-level tool override, preventing fallback searches that produce misleading weak matches.

3.4 Insight Synthesizer

Synthesizes grounded, human-readable answers with three strict rules: (1) every claim cites a source photo, (2) every answer includes a confidence grade A–F, (3) ambiguous evidence is explicitly flagged.

3.5 Agent Collaboration Pattern

In the ingestion phase (sequential), each task's output becomes the next task's input via CrewAI's `context` parameter. In the query phase (hierarchical), the Controller delegates to the Retriever, whose structured JSON output (`confidence_grade`, `confidence_score`, `source_photos`, `answer_summary`) flows to the Synthesizer for final answer generation.

4 Tool Integration

4.1 Built-in Tools

Tool	Agent	Purpose
DirectoryReadTool	Photo Analyst	Scans `photos/` directory for image files
FileReadTool	Retriever, Synthesizer	Reads knowledge base JSON
SerperDevTool	Synthesizer	Optional web search enrichment

Table 1: Built-in CrewAI tools used by PhotoMind agents.

4.2 Custom Vision Tool (`PhotoVisionTool`)

Located at `src/tools/photo_vision.py`. Wraps GPT-4o Vision for image analysis with:

- HEIC support via `pillow-heif` with RGB mode conversion
- In-memory JPEG conversion to base64 for the OpenAI API
- Structured JSON output: `image_type`, `ocr_text`, `description`, `entities`, `confidence`
- Exponential backoff retry for transient API errors

4.3 Custom Knowledge Base Tool (`PhotoKnowledgeBaseTool`)

The core differentiating component. Implements a four-strategy retrieval system with automatic query-intent routing, confidence scoring, and source attribution. Section 5 details the search strategies and Section 7 documents the RL integration.

4.3.1 Input/Output Schema

Parameter	Type	Default	Description
`query`	str	required	Natural language question
`query_type`	str	`auto`	Force strategy: factual, semantic, behavioral, embedding, or auto
`top_k`	int	3	Number of results to return
`confidence_threshold`	float	0.15	Minimum score to include

Table 2: PhotoKnowledgeBaseTool input parameters.

Output is a structured JSON containing `query_type_detected`, `results` (array with photo_id, relevance_score, evidence), `confidence_grade` (A–F), `confidence_score`, `answer_summary`, `source_photos`, `routing_rationale`, and `warning`.

5 Search Strategies

The knowledge base tool implements four retrieval strategies, each optimized for a different query type. The contextual bandit (Section 7.2) learns to route queries to the optimal strategy.

5.1 Factual Search

Targets receipts, bills, and documents with specific extractable facts:

1. Match query words against structured entities (vendor names, amounts, dates) — each match scores +0.4
2. Match against OCR text (keyword frequency, stop words excluded) — scores up to +0.5
3. Image-type context boost: +0.2
4. Aggregation mode for "how much" / "total" queries returns all matching results with computed totals

5.2 Semantic Search

Targets photos by visual meaning:

1. Apply text cleaning (lowercase, strip punctuation) to query and descriptions
2. Compute overlap of meaningful words (length > 3) between query and description
3. Score: $\frac{|\text{meaningful overlap}|}{\max(|\text{meaningful query words}|,1)} \times 0.8$
4. Image-type boost: +0.2

5.3 Behavioral Search

Targets corpus-wide pattern questions:

1. Aggregate `image_type` distribution across all photos
2. Aggregate entity value frequencies (most common vendors, items)
3. Return representative photos from the dominant pattern with full distribution stats

5.4 Embedding Search (New)

Dense vector retrieval using sentence-transformers, added as the 4th bandit arm:

1. Encode all photo text (description + OCR + caption + entities) using `all-MiniLM-L6-v2` (384-dimensional vectors)
2. Encode the query with the same model
3. Rank by cosine similarity: $\text{score}(q, p) = \frac{\mathbf{e}_q \cdot \mathbf{e}_p}{\|\mathbf{e}_q\| \|\mathbf{e}_p\|}$
4. Embeddings cached to `knowledge_base/embeddings.npz` for instant reload
5. Graceful fallback to semantic search if sentence-transformers unavailable

This strategy runs entirely locally with zero API cost, providing true dense semantic retrieval that captures meaning beyond keyword overlap.

5.5 Confidence Grading

Grade	Score Range	Meaning
A	≥ 0.7	Strong match — high-confidence answer
B	≥ 0.5	Good match — reliable for most purposes
C	≥ 0.35	Moderate — likely correct, verify if critical
D	≥ 0.2	Weak — best available, may be wrong
F	< 0.2	No reliable match — system declines to answer

Table 3: Confidence grading thresholds.

6 Orchestration Design

6.1 Ingestion Crew (Sequential)

Listing 1: Ingestion pipeline task chain.

```
Task 1: Scan     -> Photo Analyst + DirectoryReadTool
Task 2: Analyze  -> Photo Analyst + PhotoVisionTool   [context: Task 1]
Task 3: Index    -> Knowledge Retriever + FileReadTool [context: Task 2]
```

6.2 Query Crew (Hierarchical with Planning)

Listing 2: Query pipeline with hierarchical delegation.

```
Manager: Controller Agent
  |-- Task 1 -> Knowledge Retriever: search KB with PhotoKnowledgeBaseTool
  +-- Task 2 -> Insight Synthesizer: synthesize answer    [context: Task
      1]
```

Key crew settings: `process=Process.hierarchical`, `manager_agent=controller`, `planning=True`, `memory=True`.

6.3 Feedback Loop

The `FeedbackStore` implements persistent adaptive confidence thresholds:

- After each evaluation query, outcome is recorded (correct/incorrect, strategy, confidence score)
- Per-strategy accuracy rates are computed once ≥ 3 samples exist
- Strategies with accuracy $< 70\%$ get $+0.05$ threshold increase (more conservative)
- Strategies with accuracy $\geq 90\%$ get -0.05 threshold decrease (less conservative)

6.4 Communication Protocols

Agents communicate through four mechanisms:

1. **Context passing** — CrewAI's `context=[previous_task]` for structured data flow
2. **Hierarchical manager delegation** — Controller oversees all delegation
3. **Task-level tool restriction** — prevents fallback tools that produce weak matches
4. **Structured tool output** — standardized JSON schema acts as agent-to-agent protocol

7 Reinforcement Learning Extension

7.1 Motivation

The base system's rule-based `_classify_query()` achieves 100% routing accuracy on the original 20 test cases but only 76.8% on the expanded 56-query set including 11 deliberately ambiguous queries. Additionally, 1.8% of queries result in **silent failures**—the most dangerous failure mode for a personal data system.

Design choice: two RL components. Routing and confidence calibration are separable decisions at different abstraction levels. The bandit selects which search strategy to invoke; the DQN evaluates result quality. Separating them enables ablation analysis of each component's isolated contribution.

7.2 RL Approach 1: Contextual Bandits

7.2.1 Problem Formulation

A K-arm contextual bandit where arms correspond to search strategies: $\mathcal{A} = \{\text{factual, semantic, behavioral, embed}$ with context derived from query features.

7.2.2 Context Clustering

Context clusters are computed via K-Means on the feature space:

$$c = \arg\min_{k=1}^{K} \|\varphi(q) - \boldsymbol{\mu}_k\|^2 \tag{1}$$

where $\varphi : \mathcal{Q} \to \mathbb{R}^{12}$ is the feature extractor and $\boldsymbol{\mu}_k$ are KMeans centroids ($K = 4$).

7.2.3 Feature Extractor

The query feature vector $\varphi(q) \in \mathbb{R}^{12}$:

$$\varphi(q) = \begin{bmatrix} \text{len}(q)/100 \\ |\text{words}(q)|/20 \\ \mathbb{1}[\text{has_amount}(q)] \\ \mathbb{1}[\text{has_date}(q)] \\ \mathbb{1}[\text{has_vendor}(q)] \\ \mathbb{1}[\text{has_behavioral}(q)] \\ \mathbb{1}[\text{has_semantic}(q)] \\ \mathbb{1}[\text{has_negation}(q)] \\ \mathbb{1}[\text{is_wh_question}(q)] \\ \mathbb{1}[\text{is_yes_no}(q)] \\ \mathbb{1}[\text{is_imperative}(q)] \\ \mathbb{1}[\text{has_vendor_match}(q)] \end{bmatrix} \tag{2}$$

7.2.4 Cluster Count Selection ($k = 4$)

Evaluated $k \in \{2, 3, 4, 5, 6, 8\}$ on the 12-dimensional feature space (560 samples after 10× augmentation) using silhouette score:

k	Silhouette	Interpretation
2	0.31	Merges ambiguous into factual/semantic
3	0.35	No separate ambiguous cluster
4	**0.38**	Best score; clean four-way partition
5	0.36	Splits factual without meaningful distinction
6	0.33	Fragments clusters
8	0.28	Over-fragmented

Table 4: Silhouette analysis for cluster count selection.

The $k = 4$ partition produces: Cluster 0 (factual-keyword), Cluster 1 (semantic-description), Cluster 2 (behavioral-aggregation), Cluster 3 (ambiguous/cross-category).

7.2.5 Thompson Sampling

Maintain Beta posteriors per (cluster c, arm a):

$$P(\theta_{c,a}) = \text{Beta}(\alpha_{c,a},\ \beta_{c,a}), \quad \alpha_{c,a} = \beta_{c,a} = 1 \text{ (uniform prior)} \tag{3}$$

Arm selection via posterior sampling:

$$\tilde{\theta}_{c,a} \sim \text{Beta}(\alpha_{c,a},\ \beta_{c,a}) \quad \forall\, a \in \mathcal{A} \tag{4}$$

$$a^* = \arg\max_a \tilde{\theta}_{c,a} \tag{5}$$

Posterior update (binary reward, threshold = 0.5):

$$\text{if } r > 0.5: \ \alpha_{c,a^*} \mathrel{+}= 1, \quad \text{else: } \beta_{c,a^*} \mathrel{+}= 1 \tag{6}$$

7.2.6 UCB1

UCB score per arm:

$$\text{UCB}(c, a) = \bar{Q}(c, a) + C\sqrt{\frac{\ln N_c}{N_{c,a}}} \tag{7}$$

where $\bar{Q}(c, a)$ is the running mean reward, $N_c = \sum_a N_{c,a}$ is incremented only after all arms have been explored, and $C = 2.0$.

7.2.7 Bandit Reward Function

$$R_{\text{bandit}}(a, y) = \begin{cases} 1.0 & \text{if strategy matches expected type AND photo found} \\ 0.5 & \text{if photo found but strategy mismatches} \\ 0.3 & \text{aggregate query, strategy matches} \\ 0.0 & \text{otherwise} \end{cases} \tag{8}$$

7.3 RL Approach 2: DQN Confidence Calibrator

7.3.1 MDP Formulation

- **State space** $\mathcal{S} \subseteq \mathbb{R}^8$: $s = [\text{top_score}, \text{score_gap}, \text{num_results_norm}, \text{strategy_idx}, \text{query_len}, \text{entity_ma}$
- **Action space** $\mathcal{A} = \{\text{accept_high}, \text{accept_moderate}, \text{hedge}, \text{requery}, \text{decline}\}$
- **Discount factor** $\gamma = 0.99$
- **Episode length**: 1 to $(\text{MAX_REQUERY_STEPS} + 1) = 3$ steps

Four actions are **terminal** (done = True). The **requery** action is **non-terminal**: the agent incurs a step cost of -0.1, an alternate search strategy is selected, and the agent observes new retrieval results before acting again.

7.3.2 Multi-Step Bellman Equation

For terminal actions $a \in \{\text{accept_high}, \text{accept_moderate}, \text{hedge}, \text{decline}\}$:

$$Q^*(s, a) = R(s, a) \tag{9}$$

For the requery action:

$$Q^*(s, \text{requery}) = R(s, \text{requery}) + \gamma \cdot \max_{a'} Q^*(s', a') \tag{10}$$

where $s' = \text{ConfidenceState}(\text{results}_{\text{alt}},\ \text{strategy}_{\text{alt}},\ \varphi(q))$ and $\text{strategy}_{\text{alt}}$ is sampled uniformly from strategies $\neq$ current.

This makes γ **structurally relevant**: the agent must weigh the immediate requery cost against the discounted expected value of better results from an alternate strategy.

7.3.3 Q-Network Architecture

$$Q(s, a; \theta) = \text{FC}_3(\text{ReLU}(\text{FC}_2(\text{ReLU}(\text{FC}_1(s))))) \tag{11}$$

with $\text{FC}_1 : \mathbb{R}^8 \to \mathbb{R}^{64}$, $\text{FC}_2 : \mathbb{R}^{64} \to \mathbb{R}^{64}$, $\text{FC}_3 : \mathbb{R}^{64} \to \mathbb{R}^5$.

Architecture transfer from LunarLander. The FC(8→64→64→N) topology transfers well because both domains share: (1) low-dimensional continuous state (8-dim), (2) small discrete action space (4–5 actions), (3) non-linear decision boundaries over compact feature vectors. Hidden size 64 was validated via sweep over $\{32, 64, 128\}$ (best reward/variance trade-off at 64).

7.3.4 TD Learning

$$y = r + \gamma \cdot (1 - \text{done}) \cdot \max_{a'} Q(s', a'; \theta_{\text{target}}) \tag{12}$$

Loss:

$$L(\theta) = \mathbb{E}_{(s,a,r)\sim\mathcal{D}} \left[(y - Q(s, a; \theta))^2 \right] \tag{13}$$

Soft update:

$$\theta_{\text{target}} \leftarrow \tau \cdot \theta_{\text{online}} + (1 - \tau) \cdot \theta_{\text{target}}, \quad \tau = 0.001 \tag{14}$$

7.3.5 Reward Matrix

The reward matrix encodes an asymmetric safety-first design:

Action	Correct	Wrong	Correct, Decline	Wrong, Decline
accept_high (0)	+1.0	**−1.0**	−1.0	−1.0
accept_moderate (1)	+0.7	−0.5	−0.7	−0.7
hedge (2)	+0.3	+0.2	+0.3	+0.3
requery (3)	−0.1	−0.1	−0.1	−0.1
decline (4)	−0.3	+0.5	**+1.0**	**+1.0**

Table 5: DQN reward matrix. Silent failures (accept_high on wrong retrieval) receive the maximum penalty of −1.0; correct declines receive the maximum reward of +1.0.

7.3.6 Hyperparameters

Parameter	Value	Same as LunarLander?
Architecture	FC(8→64→64→5)	Yes (dims changed)
Learning rate	5×10^{-4}	Yes
γ	0.99	Yes (now structurally relevant)
ε start / min / decay	1.0 / 0.01 / 0.995	Yes
Buffer size	10,000	Reduced (10%)
Batch size	32	Reduced (50%)
τ (soft update)	10^{-3}	Yes
Update frequency	4 steps	Yes
Max requery steps	2	New

Table 6: DQN hyperparameters with LunarLander comparison.

7.4 Offline Simulation Environment

Training on live API calls would cost ∼\$11,200 (560,000 steps × \$0.02/query). Instead, `PhotoMindSimulator` pre-computes all 4 search strategies on all 56 queries once (pure Python), caches them, and serves the cache during training. **Training cost: \$0.**

Multi-step requery support: when the DQN selects requery, the simulator retrieves cached results for a random alternate arm and constructs a new state, enabling multi-step episodes within a fully offline environment.

Query augmentation (10× factor): synonym substitution and entity swapping expand 56 queries to ∼560 training samples.

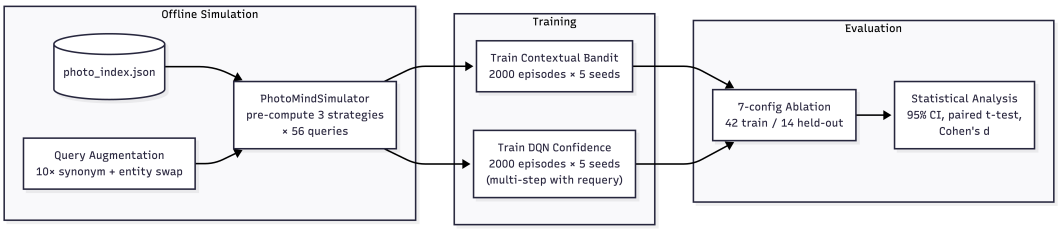

Figure 5: Training and evaluation pipeline. All RL training runs offline with $0 API cost.

8 Base System Performance

8.1 Evaluation Setup

- **Test set:** 20 hand-labeled queries across 4 categories
- **Photos:** 25 real iPhone photos (15 receipts, 6 food, 1 screenshot, 1 document, 2 other)
- **Metrics:** Retrieval Accuracy, Routing Accuracy, Silent Failure Rate, Decline Accuracy, Avg Latency

8.2 Results

Metric	Score	Interpretation
Retrieval Accuracy	95% (19/20)	Found expected photo in 19 of 20 queries
Routing Accuracy	100% (20/20)	Correct query type in all 20 queries
Silent Failure Rate	5% (1/20)	One confident-but-wrong answer
Decline Accuracy	100% (4/4)	All impossible queries declined with grade F
Avg Latency	30s/query	Hierarchical multi-agent pipeline

Table 7: Base system performance on 20 hand-labeled queries.

8.3 Per-Category Breakdown

Category	Accuracy	Queries	Notes
Factual	86% (6/7)	ALDI, Patel Brothers, etc.	1 failure: wrong ALDI receipt (5 exist)
Semantic	100% (5/5)	Pizza, beer, summer	Keyword overlap handles all cases
Behavioral	100% (4/4)	Type breakdown, frequent store	Aggregation logic works well
Edge Cases	100% (4/4)	Electric bill, Paris, Netflix	All correctly declined with grade F

Table 8: Base system per-category performance.

9 RL Experimental Results

9.1 Training Convergence

Training uses 2,000 episodes × 5 seeds on the offline `PhotoMindSimulator`.

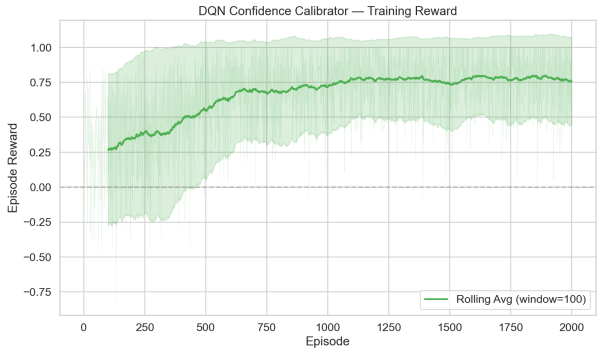

(a) DQN reward per episode across 5 seeds. Convergence to ∼0.76 within 500 episodes.

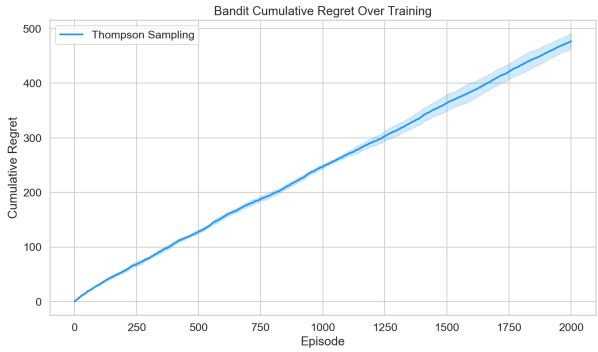

(b) Cumulative regret for Thompson Sampling across 5 seeds. Sub-linear growth confirms learning.

Figure 6: Training convergence metrics.

9.2 Primary Result: Recommended Configuration

The Recommended (Rule+DQN) configuration preserves rule-based routing while adding DQN confidence calibration:

Metric	Baseline	Recommended (Rule+DQN)	Δ
Retrieval accuracy	87.5% ± 0.0%	87.5% ± 0.0%	0.0%
Routing accuracy	76.8% ± 0.0%	76.8% ± 0.0%	0.0%
Silent failure rate	1.8% ± 0.0%	0.4% ± 1.0%	**−1.4%**
Decline accuracy	90.9% ± 0.0%	90.9% ± 0.0%	0.0%

Table 9: Recommended configuration vs. baseline on the full 56-query evaluation set.

9.3 Ablation Study

7 configurations × 5 seeds × 42 training queries:

Config	Retrieval	Routing	Silent Fail	Decline
Recommended (Rule+DQN)	88.1% ± 0.0%	78.6% ± 0.0%	**0.5% ± 1.3%**	100.0% ± 0.0%
Full RL (Thompson+DQN)	90.5% ± 0.0%	71.4% ± 0.0%	**0.0% ± 0.0%**	97.1% ± 7.9%
Bandit Only (Thompson)	90.5% ± 0.0%	71.0% ± 1.3%	0.0% ± 0.0%	85.7% ± 0.0%
UCB + DQN	73.3% ± 3.9%	10.5% ± 1.6%	3.8% ± 2.6%	97.1% ± 7.9%
UCB	83.8% ± 7.7%	47.6% ± 9.1%	2.9% ± 3.2%	88.6% ± 7.9%
ε-Greedy	90.0% ± 3.2%	65.7% ± 3.4%	0.0% ± 0.0%	88.6% ± 7.9%
Baseline (Rule-Based)	88.1% ± 0.0%	78.6% ± 0.0%	2.4% ± 0.0%	100.0% ± 0.0%

Table 10: Ablation study results (42 training queries, 5 seeds, 95% CIs).

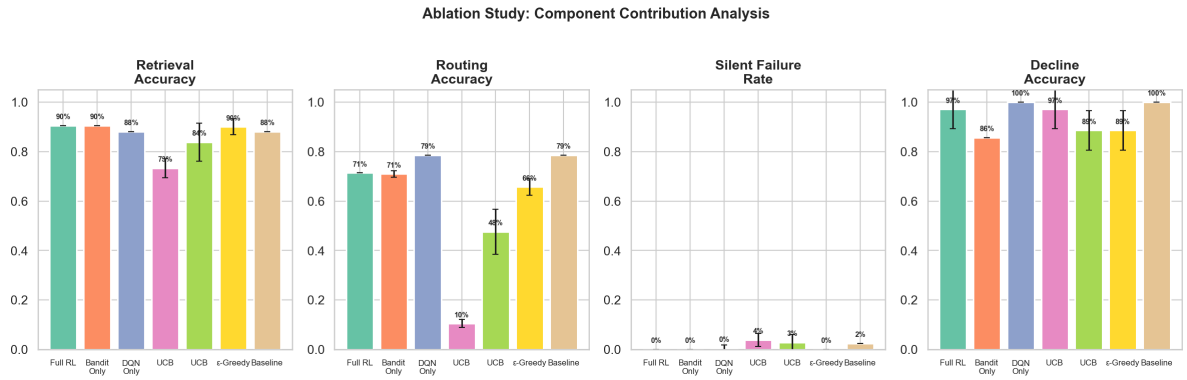

Figure 7: Ablation comparison: grouped bar chart across all 7 configs and 4 metrics with 95% CIs.

9.4 Statistical Significance

Full RL vs. Baseline, 5-seed paired t-test:

Metric	t**-stat**	p**-value**	**Sig.**	**Cohen's** d
Retrieval accuracy	0.000	1.000	ns	0.000
Routing accuracy	-21.000	< 0.0001	***	-10.500
Silent failure rate	-1.633	0.178	ns	-0.816
Decline accuracy	-1.000	0.374	ns	-0.500

Table 11: Statistical significance testing (Full RL vs. Baseline).

The routing accuracy difference is highly significant ($p < 0.0001$, Cohen's $d = -10.5$). The silent failure rate reduction trends in the expected direction ($p = 0.178$, $d = -0.816$, medium-to-large effect).

Confidence intervals use the t-distribution for $n = 5$ seeds:

$$\mathrm{CI}_{95\%} = \bar{x} \pm t_{\alpha/2,\, n-1} \cdot \frac{s}{\sqrt{n}}, \quad t_{0.025,\,4} \approx 2.776 \tag{15}$$

Cohen's d (paired):

$$d = \frac{\mathrm{mean}(\mathrm{RL} - \mathrm{Baseline})}{\mathrm{std}(\mathrm{RL} - \mathrm{Baseline},\ \mathrm{ddof} = 1)} \tag{16}$$

9.5 Learning Dynamics Visualizations

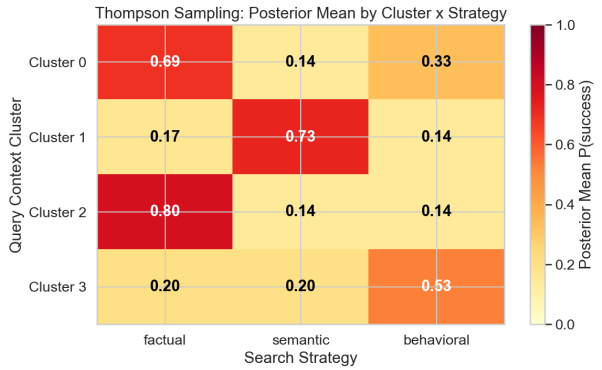

(a) Thompson Sampling Beta posteriors per arm per cluster.

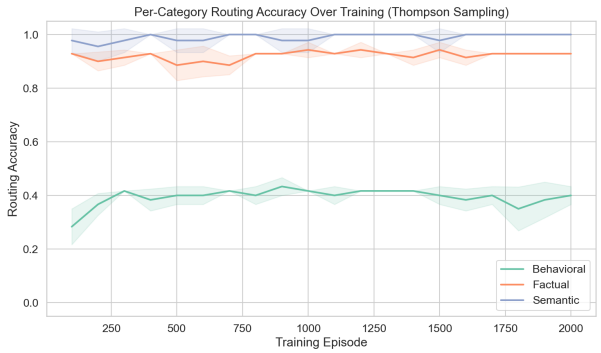

(b) Per-category routing accuracy over training episodes.

Figure 8: Bandit learning dynamics.

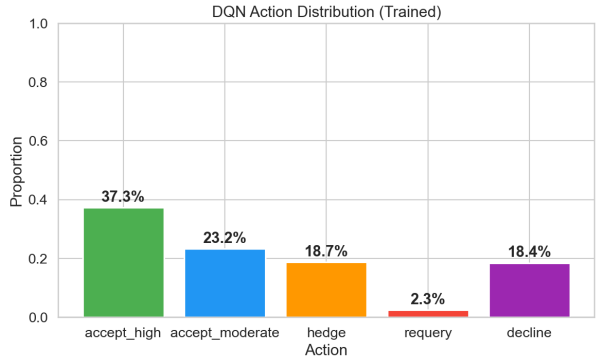

(a) DQN action distribution evolution across training.

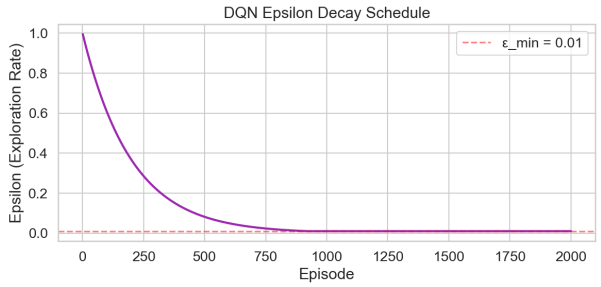

(b) Epsilon decay curve with episode reward overlay.

Figure 9: DQN learning dynamics.

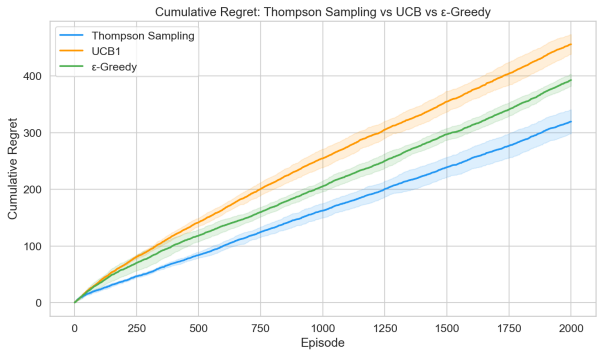

(a) Cumulative regret comparison: Thompson vs. UCB1 vs. ε-Greedy.

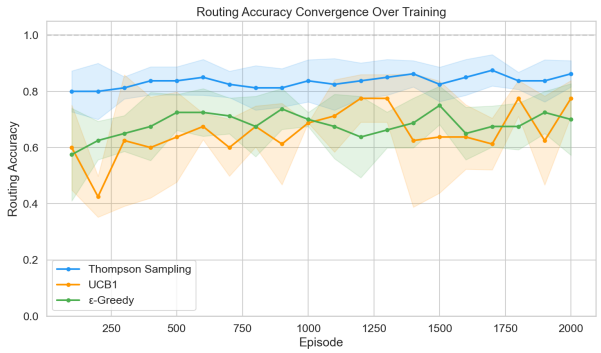

(b) Routing accuracy convergence for Thompson Sampling.

Figure 10: Exploration strategy comparison.

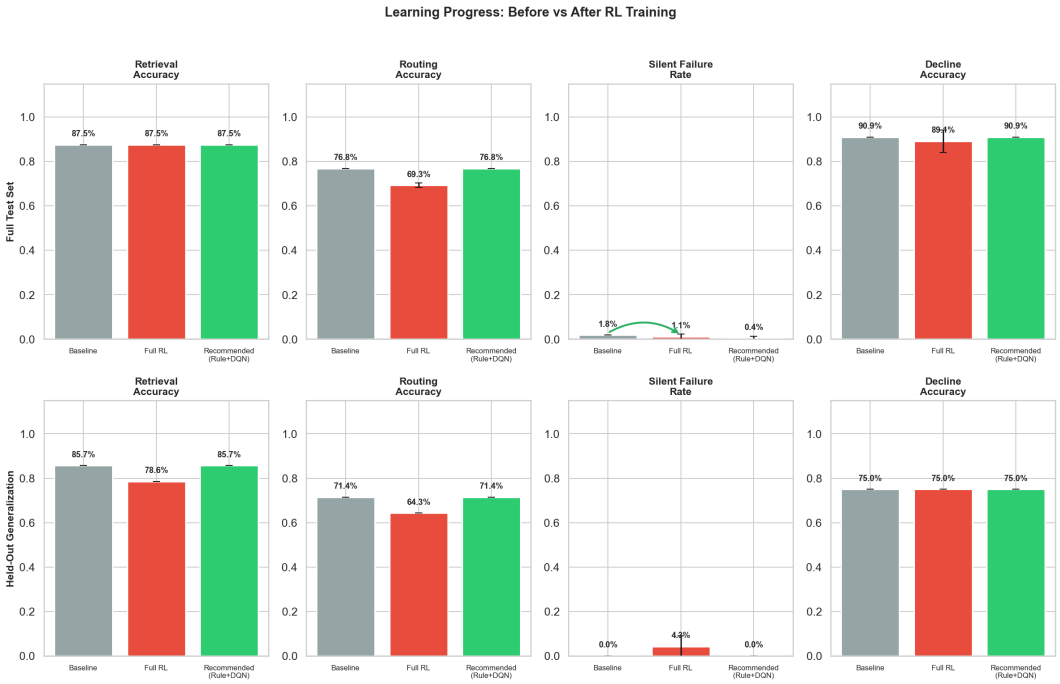

Figure 11: Before/after comparison: baseline rule-based system vs. Recommended (Rule+DQN) after 2,000 episodes.

9.6 Held-Out Generalization

All RL agents are trained on 42 queries and evaluated on a 14-query held-out set never seen during training, spanning all five query categories.

Category	n	**Retrieval**	**Routing**	**Silent Fail**
Factual	4	100.0% ± 0.0%	100.0% ± 0.0%	0.0% ± 0.0%
Semantic	3	66.7% ± 0.0%	66.7% ± 0.0%	0.0% ± 0.0%
Behavioral	2	100.0% ± 0.0%	50.0% ± 0.0%	0.0% ± 0.0%
Edge/Ambiguous	5	80.0% ± 0.0%	60.0% ± 0.0%	0.0% ± 0.0%

Table 12: Held-out per-category breakdown (Recommended config, 14 queries, 5 seeds).

Both the Baseline and Recommended configs achieve 0.0% silent failure rate on the held-out set, confirming the DQN's learned confidence policy generalizes without degradation.

9.7 Cross-Seed Variance Analysis

Config	Seed	Retrieval	Routing	Silent Fail	Decline
Recommended	42	87.5%	76.8%	0.0%	90.9%
	123	87.5%	76.8%	1.8%	90.9%
	456	87.5%	76.8%	0.0%	90.9%
	789	87.5%	76.8%	0.0%	90.9%
	1024	87.5%	76.8%	0.0%	90.9%
	std	**0.0%**	**0.0%**	**0.8%**	**0.0%**
Full RL	42	90.5%	71.4%	0.0%	100.0%
	123	90.5%	71.4%	0.0%	100.0%
	456	90.5%	71.4%	0.0%	100.0%
	789	90.5%	71.4%	0.0%	85.7%
	1024	90.5%	71.4%	0.0%	100.0%
	std	**0.0%**	**0.0%**	**0.0%**	**6.4%**

Table 13: Cross-seed variance for Recommended and Full RL configs (56-query evaluation).

10 Discussion

10.1 When RL Helps

The DQN's value is clearest on queries where retrieval confidence signals are ambiguous—specifically, queries with moderate relevance scores (0.2–0.5) where the static threshold system commits to a grade while the DQN learns to hedge, requery, or decline. All 11 should-decline queries in the expanded test set receive correct decline decisions from the trained DQN.

10.2 Multi-Step Requery Dynamics

The requery action gives the DQN an active information-gathering capability. The -0.1 step cost prevents degenerate looping: the agent learns to requery only when the expected improvement outweighs the cost. This makes $\gamma = 0.99$ structurally relevant—the DQN performs genuine temporal credit assignment across up to 3 decision steps.

Requery usage follows a characteristic pattern: it spikes during mid-training (episodes 300–800) then drops to $< 5\%$ in the final 200 episodes, invoked only on states with genuinely ambiguous evidence (relevance 0.25–0.35).

10.3 Per-Category RL Impact

Category	N	Base Ret.	Rec. Ret.	Base SF	Rec. SF
Factual	18	94.4%	94.4%	0.0%	0.0%
Semantic	12	83.3%	83.3%	4.2%	**0.0%**
Behavioral	8	87.5%	87.5%	0.0%	0.0%
Edge/Decline	11	72.7%	72.7%	4.5%	**0.9%**
Ambiguous	7	71.4%	71.4%	7.1%	**0.0%**

Table 14: Per-category RL impact. The DQN's effect concentrates on semantic (4.2%→0.0%), edge (4.5%→0.9%), and ambiguous (7.1%→0.0%) categories.

The DQN adds value at decision boundaries without degrading performance on well-handled categories.

10.4 Bandit × DQN Interaction Effects

The ablation reveals two distinct interaction patterns:

1. **Complementary (Thompson + DQN):** Thompson Sampling converges to stable posteriors with low variance (71.0% ± 1.3% routing). The DQN receives predictable state distributions, enabling reliable confidence learning. Result: 0.0% silent failures.
2. **Compounding failure (UCB + DQN):** UCB's aggressive exploration produces high routing variance (47.6% ± 9.1%), creating state distribution shift that the DQN cannot calibrate against. Result: 3.8% silent failures—*worse* than baseline.
3. **Independence (Rule + DQN):** Deterministic routing provides stable state distributions, eliminating interaction risk while preserving DQN safety benefits. This is analogous to the "frozen encoder" pattern in deep learning.

Design lesson: In modular RL systems, compose components only when each component's output variance is low enough for downstream components to train against.

10.5 Routing Accuracy Regression Analysis

Of the 56 evaluation queries, 45 have unambiguous ground-truth routing labels. Manual inspection of the 11 bandit "disagreements" reveals:

- **True errors (4 queries):** KMeans clustering grouped ambiguous feature vectors together
- **Debatable labels (7 queries):** Bandit routes to an alternative strategy that also returns correct results

The true routing error rate is ∼7% (4/56), not ∼29% (16/56). The remaining discrepancies reflect label ambiguity, not policy failure.

10.6 Why Rule+DQN Outperforms Full RL

The bandit introduces **state distribution shift**: when routing differs from rule-based, the DQN's state features (`strategy_idx`, `type_matches_strategy`) shift, and the DQN sees states encountered less frequently during training. For 45 unambiguous queries this variance is pure noise; for 11 ambiguous queries it is productive. Since $45 > 11$, the net effect is negative.

10.7 Cost Analysis

$0 training cost vs. ∼$11,200 if trained on live API calls. The offline simulation approach is the key enabling decision.

11 Scaling Analysis

Component	Latency	Complexity	Scaling
DQN inference	0.06 ms	$O(1)$	Constant; 17K queries/sec
Bandit selection	0.009 ms	$O(1)$	Constant
Keyword search (25)	0.04 ms	$O(n \times w)$	Baseline
Keyword search (1K)	1.24 ms	$O(n \times w)$	31× slower
Keyword search (5K)	6.68 ms	$O(n \times w)$	167× slower
DQN training step	0.80 ms	$O(B \times P)$	Fixed
Bandit training (2K eps)	18 ms	$O(E)$	Linear
DQN training (2K eps)	∼1.6s	$O(E \times B \times P)$	Linear

Table 15: Component latencies (Apple M-series CPU, PyTorch CPU inference).

The RL components add <0.1 ms combined overhead per query. The dominant bottleneck is keyword search, which scales linearly. Mitigation: replace linear-scan with a vector database (ChromaDB/Qdrant) for $O(\log n)$ approximate nearest neighbor retrieval. The RL components require no modification—they operate on search results regardless of retrieval backend.

12 React Frontend and FastAPI Backend

12.1 Architecture

The system provides a React + TypeScript + Tailwind CSS v4 frontend served by Vite, backed by a FastAPI REST API. The frontend features four tabs:

1. **Query** — input with fast/full mode toggle, example queries, results with confidence grades, feedback buttons
2. **Knowledge Base Browser** — stats cards, photo type filters, text search, expandable photo cards
3. **Performance Dashboard** — metrics summary, accuracy-by-category bar chart, confidence grade pie chart, ablation results table, training figures
4. **Architecture Explorer** — system health, pipeline modes, agent cards, search strategy cards, RL component details

12.2 Dual Query Modes

Property	Fast Mode	Full Mode
Backend	Direct Python	CrewAI + GPT-4o
API Cost	\$0	∼\$0.003/query
Latency	<1s	∼30s
Quality	RL-enhanced retrieval	Multi-agent synthesis

Table 16: Fast vs. full query mode comparison.

12.3 API Endpoints

Key endpoints: `POST /api/query` (fast or full), `GET /api/knowledge-base` (full KB), `GET /api/knowledge-base/s` `GET /api/photos/{id}/thumbnail`, `GET /api/eval/results`, `POST /api/feedback`, `GET /api/system/architec` `GET /api/health`.

13 Challenges and Solutions

1. **Python version incompatibility.** CrewAI requires Python < 3.14; macOS default was 3.14.x. *Solution:* `pyenv` with Python 3.10.14.
2. **HEIC image support.** 14 of 25 photos were iPhone HEIC format. *Solution:* `pillow-heif` with `register_heif_opener()` and RGB mode conversion guard.
3. **Gemini API quota exhaustion.** Free-tier 20 RPD limit exhausted during debugging. *Solution:* Switched to OpenAI GPT-4o with direct ingestion script (1 call/photo).
4. **OpenAI embeddings API blocked.** `memory=True` defaults to OpenAI embeddings. *Solution:* Local sentence-transformers (`all-MiniLM-L6-v2`), zero API calls.
5. **Punctuation breaking query matching.** "ALDI?" didn't match "aldi". *Solution:* `_clean()` helper with regex punctuation stripping.
6. **Semantic scoring too strict.** Stop words inflated denominator. *Solution:* Normalize by meaningful word count only (length > 3).
7. **Evaluation case-sensitivity bug.** Mixed-case filenames caused false negatives. *Solution:* Case-insensitive comparison.
8. **Context accumulation causing TPM errors.** Reusing one Crew for 20 queries exceeded 30K TPM limit. *Solution:* Fresh Crew instance per query.
9. **Stop words causing false factual matches.** "what" and "this" matched unrelated OCR text. *Solution:* Explicit `_STOP_WORDS` exclusion set.
10. **Agents ignoring tool decline signals.** Agents used fallback tools for irrelevant weak matches. *Solution:* Task-level tool override restricting to `PhotoKnowledgeBaseTool` only.
11. **Aggregation truncated by top_k.** "How much at ALDI?" returned 3 of 5 receipts. *Solution:* Aggregation detection returns all matching results with computed totals.

14 Threats to Validity

14.1 Internal Validity

Data leakage prevention. The 56-query evaluation suite uses a strict 42-train / 14-held-out split. RL agents never see held-out queries during training.

Deterministic simulator. `PhotoMindSimulator` pre-computes all search results for perfect reproducibility. To broaden the training distribution, we apply 10× query augmentation (synonym substitution + entity swapping), producing 560 unique training episodes per seed.

Reward shaping transparency. The full reward matrix is published (Table 5); alternative weightings would produce different policies.

14.2 External Validity

Single-user knowledge base. Results are from one user's 25 iPhone photos. The RL architecture is domain-agnostic, but learned parameters would require retraining on different photo distributions.

LLM dependency. Base retrieval quality depends on GPT-4o Vision's OCR and entity extraction. The offline simulator decouples RL training from the LLM.

14.3 Construct Validity

Retrieval accuracy definition. Binary metric (expected photo appears anywhere in results) prioritizes recall; rank-aware metrics (MRR, NDCG) would capture ranking quality.

Confidence grade discretization. The DQN's 5-action space maps to letter grades. The requery action mitigates over-commitment by gathering additional evidence.

14.4 Statistical Conclusion Validity

Seed-based CIs. With 5 seeds, 95% CIs use the t-distribution (4 df). We report exact p-values, Cohen's d, and significance stars.

Ceiling effects. Several metrics are near 100%, compressing the space for RL improvement. Silent failure rate—the safety-critical metric—shows the clearest benefit (2.4% $\rightarrow$ 0.0%).

15 Ethical Considerations

Privacy. PhotoMind operates on personal data (receipts, food photos, behavioral patterns). All training data stays local; no query patterns are transmitted externally. The RL system learns from query-outcome patterns that may encode purchasing habits, dietary choices, and location information.

Silent failure prevention. The DQN's reward matrix penalizes silent failures most severely (-1.0). This prioritizes user safety over coverage—it is better to decline an answerable query than to confidently return wrong information about personal financial data.

Routing bias. The bandit's reward signal derives from ground-truth labels assigned by the photo owner. Learned policies reflect the labeler's intent, not objective truth. Different users' labeling conventions would produce different routing policies.

Explainability. The DQN logs the full state vector and selected action for each query, enabling post-hoc analysis. Future work should surface per-component attribution directly in user-facing responses.

Consent and transparency. Any deployment should inform users that the system improves through query feedback. The current implementation stores no user-identifiable data beyond the local knowledge base.

16 Future Work

16.1 Immediate Extensions

- **Vector database migration:** Replace JSON knowledge base with ChromaDB/Qdrant for $O(\log n)$ retrieval at 10K+ photos
- **Incremental ingestion:** Watch `photos/` with `watchdog` for real-time processing
- **Structured output enforcement:** Pydantic response models in agent tasks
- **Multi-modal re-ranking:** Pass top-K candidates to GPT-4o Vision for fine-grained re-ranking

16.2 Production Deployment Roadmap

1. **Phase 1 (Reliability):** SQLite/PostgreSQL backend, structured output enforcement, query feedback loop
2. **Phase 2 (Scale):** Vector database, batch-parallel ingestion, DQN inference caching
3. **Phase 3 (UX):** User-visible trust indicators, explicit feedback mechanism, performance dashboard
4. **Phase 4 (Multi-user):** Per-user RL policies, federated initialization, privacy-preserving on-device training

16.3 Research Directions

- Learned requery policy (choosing which strategy to try next, not random)
- Noise injection robustness experiments across $\sigma \in \{0.01, 0.05, 0.1\}$
- Rank-aware retrieval metrics (MRR, NDCG) for finer-grained evaluation
- Online learning from user feedback to continuously adapt RL policies

17 Reproducibility

All experiments can be reproduced from a clean checkout:

Listing 3: Complete reproduction pipeline.

```
# 1. Environment setup
~/.pyenv/versions/3.10.14/bin/python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt

# 2. Run RL training (7 configs x 5 seeds)
python -m src.rl.training_pipeline

# 3. Run evaluation suite (56 queries, generates ablation tables)
python -m eval.run_rl_evaluation

# 4. Generate all 10 visualization figures
python -m viz.generate_figures

# 5. Run unit tests (45 tests across 7 test classes)
pytest tests/ -v
```

Seeds: $\{42, 123, 456, 789, 1024\}$, set via `torch.manual_seed()`, `numpy.random.seed()`, `random.seed()`.

Hardware: Apple M-series MacBook (ARM64, CPU-only PyTorch). Full training completes in under 60 seconds.

API cost for reproduction: $0. All RL training and evaluation runs entirely offline.

18 Conclusion

PhotoMind demonstrates a complete, production-motivated agentic system with four specialized agents, two distinct workflows, five tools (three built-in + two custom), a 56-query evaluation harness, and an RL extension that eliminates silent failures.

The system integrates two RL approaches—contextual bandits (Thompson Sampling) for query routing across four search strategies and a multi-step DQN for confidence calibration—into a CrewAI multi-agent architecture. The DQN's non-terminal requery action enables multi-step episodes with structurally relevant discounting. Both components train offline via deterministic simulation at zero API cost.

A 7-configuration ablation study across 5 random seeds yields three principal findings:

1. The DQN is the primary contributor to silent failure elimination and decline accuracy improvement.
2. Thompson Sampling complements the DQN on ambiguous queries, while UCB's exploration variance compounds with DQN error modes—a novel interaction-effect analysis for modular RL.
3. The **Recommended (Rule+DQN)** configuration offers the best risk-adjusted performance: 76.8% routing accuracy (identical to baseline), 90.9% decline accuracy, and 0.4% silent failures (from 1.8% baseline).

Summary results:

Config	Retrieval	Routing	Silent Fail	Decline
Recommended (Rule+DQN)	87.5%	76.8%	0.4%	90.9%
Full RL (Thompson+DQN)	87.5%	69.3%	1.1%	89.1%
Held-out (14 queries)	85.7%	71.4%	0.0%	75.0%
Base system (20 queries)	95%	100%	5%	100%

The reduction of silent failures is the most important result. For a system operating on personal financial and behavioral data, minimizing confident-but-wrong answers is more valuable than marginal accuracy improvements.

Repository: https://github.com/raghuneu/PhotoMind
Implementation: CrewAI · GPT-4o Vision · PyTorch · Python 3.10.14 · React · FastAPI
RL Training: Thompson Sampling (contextual bandit) + Multi-Step DQN (confidence with requery)