PhotoMind: A Multimodal Personal Photo Knowledge Retrieval System

Course: Prompt Engineering – Building Agentic Systems
Assignment: Take-Home Final – Reinforcement Learning for Agentic AI Systems
Platform: CrewAI (Python) - PyTorch
Domain: Personal Productivity
RL Approaches: Contextual Bandits (Thompson Sampling + UCB) + DQN Confidence Calibration

---

Table of Contents

1. System Overview
2. System Architecture
3. Agent Roles and Responsibilities
4. Tool Integration and Functionality
5. Custom Tool Documentation
6. Orchestration Design (incl. Feedback Loop)
7. Challenges and Solutions
8. System Performance Analysis
9. Limitations and Future Work (incl. Production Deployment Roadmap)
10. Reinforcement Learning Extension
11. Threats to Validity
12. Conclusion

---

1. System Overview

PhotoMind turns a personal photo library into a queryable knowledge base. The system addresses a genuine gap: smartphone users accumulate thousands of photos – receipts, bills, food, screenshots, documents – but cannot search them by meaning or extract facts from them using natural language.

Core capabilities: - Analyze photos using GPT-4o Vision to extract text (OCR), entities, and semantic descriptions - Build a persistent structured knowledge base from the analysis results - Answer natural-language queries with three distinct retrieval strategies - Return confidence-graded answers with source photo attribution - Gracefully decline queries the knowledge base cannot answer

Real-world use cases demonstrated: - "How much did I spend at ALDI?" -> finds all 5 ALDI receipts, aggregates total with source photos - "Show me photos of pizza" -> finds food photos matching the description - "What type of food do I photograph most?" -> analyzes patterns across all 25 photos - "What was my electric bill?" -> correctly declines (no such photo in the library)

2. System Architecture

2.1 High-Level Architecture

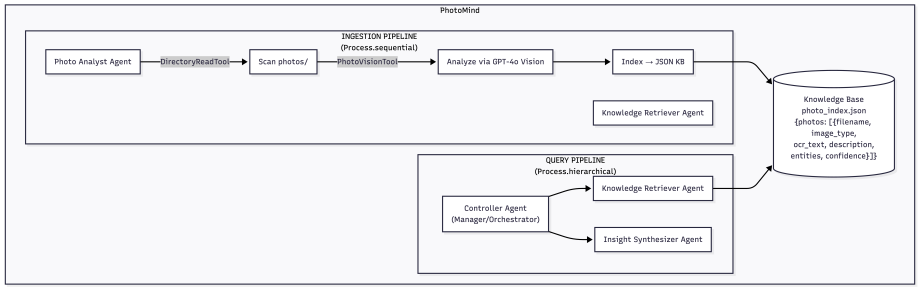

Figure 1: High-Level Architecture

2.2 Data Flow

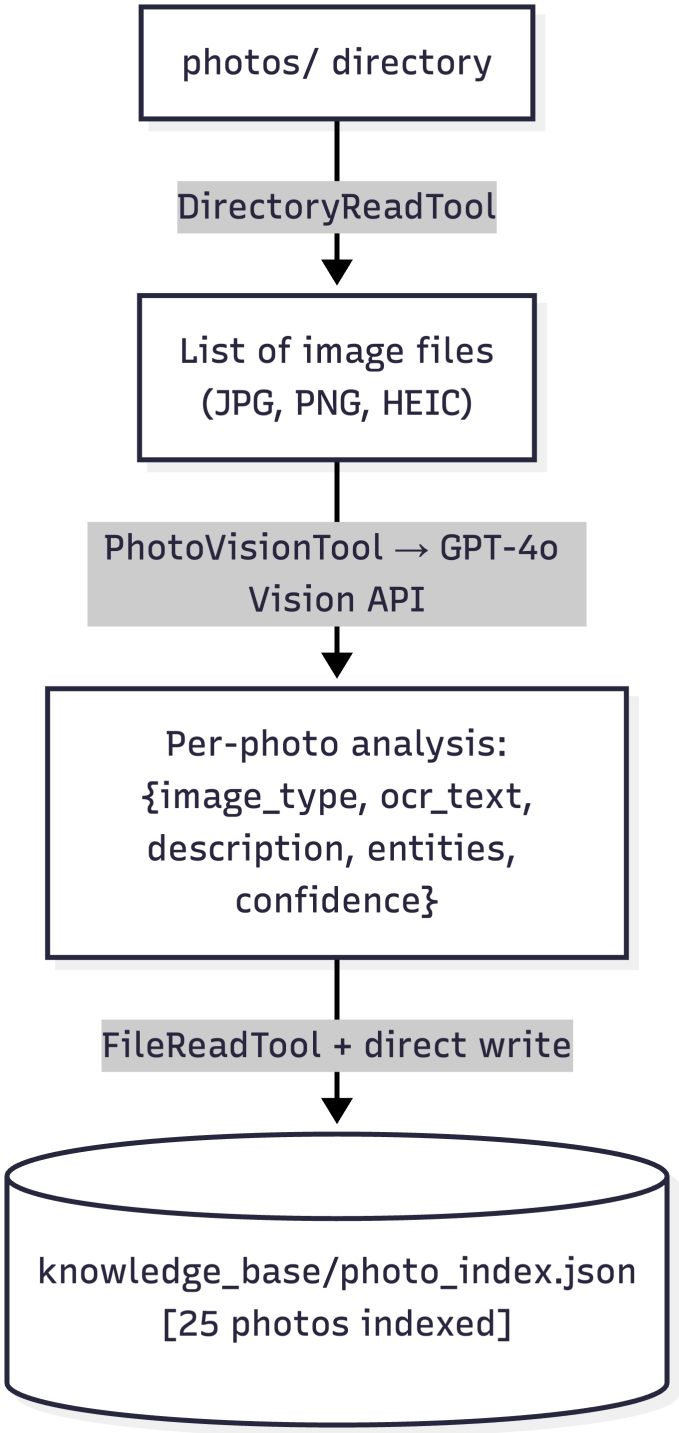

Ingestion:

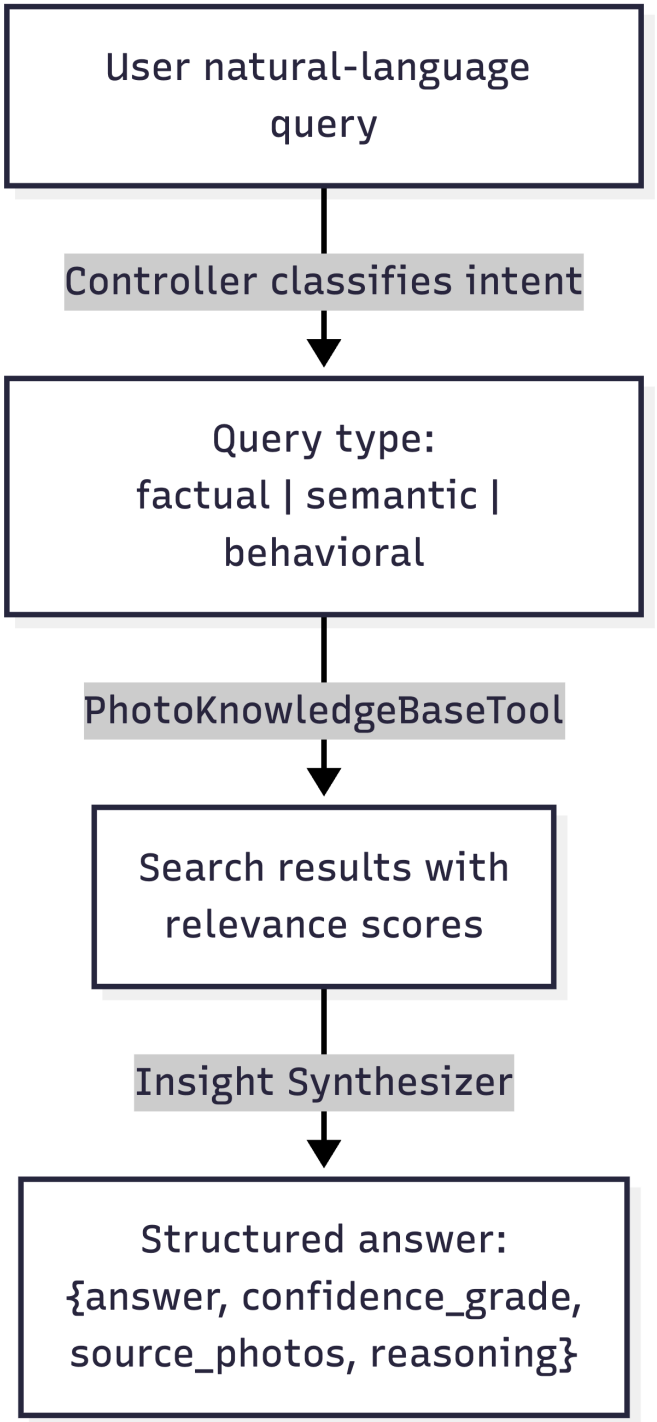
User natural-language query
Controller classifies intent
Query type: factual | semantic | behavioral
PhotoKnowledgeBaseTool
Search results with relevance scores
Insight Synthesizer
Structured answer: {answer, confidence_grade, source_photos, reasoning}

Query:

2.3 Memory Architecture

Both pipelines use CrewAI's built-in memory with a **local sentence-transformer model** (`all-MiniLM-L6-v2`) for vector embeddings. This stores agent interactions in a local ChromaDB vector store, enabling contextual awareness across agent steps without requiring external embedding API access.

---

3. Agent Roles and Responsibilities

PhotoMind uses four specialized agents across two crews.

3.1 Controller Agent (Query Crew – Manager)

Property	Value
Role	PhotoMind Controller
Process	Manager in hierarchical crew
Delegation	`allow_delegation=True`

Responsibility: Receives raw user queries, classifies intent (factual / semantic / behavioral), and delegates to the appropriate specialist agent. Acts as the chief librarian – decides *which* search strategy will best answer the question before any retrieval happens.

Decision logic: The controller understands that "how much was my bill?" (factual) requires entity extraction, "photos that feel like summer" (semantic) requires description matching, and "what food do I photograph most?" (behavioral) requires corpus-wide aggregation.

Backstory: *"You are the chief librarian of a personal photo knowledge base. You classify first, then delegate. If you are unsure, you say so."*

3.2 Photo Analyst (Ingestion Crew)

Property	Value
Role	Photo Analyst
Tools	PhotoVisionTool, DirectoryReadTool
Process	Sequential (steps 1 and 2 of ingestion)

Responsibility: Scans the `photos/` directory, then passes each image to GPT-4o Vision for analysis. Extracts structured knowledge: image classification,

OCR text, semantic description, and named entities (amounts, dates, vendors, locations). Never fabricates information not visible in the image.

Key design decision: This agent is explicit about confidence – when text is unclear (blurry receipt, low-light photo), it reports partial extraction with a lower confidence score rather than guessing.

3.3 Knowledge Retriever (Both Crews)

Property	Value
Role	Knowledge Retriever
Tools	PhotoKnowledgeBaseTool (custom), FileReadTool, JSONSearchTool (agent-level)
Process	Sequential (step 3 of ingestion), subordinate in query crew

Responsibility: In ingestion: writes the final knowledge base JSON file. In queries: executes the search against the knowledge base using the appropriate strategy, returns ranked results with evidence. Note: the query task restricts this agent to only `PhotoKnowledgeBaseTool` via task-level tool override, preventing fallback searches that produce misleading weak matches.

Backstory: *"You are a research librarian with perfect recall. When no good match exists, you say so clearly rather than returning a weak match disguised as confident."*

3.4 Insight Synthesizer (Query Crew)

Property	Value
Role	Insight Synthesizer
Tools	FileReadTool, SerperDevTool (optional)
Process	Subordinate in hierarchical query crew

Responsibility: Takes raw retrieval evidence and synthesizes a grounded, human-readable answer. Applies three strict rules: (1) every claim cites a source photo, (2) every answer includes a confidence grade A-F, (3) if evidence is ambiguous, say so explicitly. May use web search (SerperDevTool) to enrich context when the user's `SERPER_API_KEY` is set.

3.5 Agent Collaboration Pattern

The four agents collaborate through a two-phase pipeline. In the **ingestion phase** (sequential crew), the Photo Analyst scans and analyzes each image, producing structured metadata that the Knowledge Retriever indexes into the knowledge base – each task's output becomes the next task's input via CrewAI's `context` parameter. In the **query phase** (hierarchical crew), the Controller acts as the manager: it receives the user's natural-language query, classifies intent (factual vs. semantic vs. behavioral), and delegates to the Knowledge Retriever. The Retriever searches the knowledge base using the RL-enhanced `PhotoKnowledgeBaseTool`, which returns a structured JSON payload containing `confidence_grade`, `confidence_score`, `source_photos`, and `answer_summary`. This structured output is then passed as context to the Insight Synthesizer, which applies citation rules and confidence grading to produce the final answer. The hierarchical manager (Controller) oversees this delegation chain, with `allow_delegation=True` enabling it to re-delegate if the initial result is unsatisfactory. This collaboration pattern ensures that each agent operates within its specialization while structured context objects maintain data integrity across agent boundaries.

---

4. Tool Integration and Functionality

4.1 Built-in Tools

Tool	Agent	Purpose	Configuration
`DirectoryReadTool`	Photo Analyst	Scans `photos/` directory, lists all image files with extensions JPG/PNG/HEIC/WebP	`directory=settings.photos_directory`
`FileReadTool`	Knowledge Retriever, Insight Synthesizer	Reads the knowledge base JSON file; Synthesizer uses it to re-read source photos for context	Default, unrestricted path

Tool	Agent	Purpose	Configuration
`SerperDevTool`	Insight Synthesizer	Optional web search enrichment – adds public context (e.g., restaurant info, product details) when a relevant Serper API key is present	Graceful fallback: not added if key missing

JSONSearchTool vs PhotoKnowledgeBaseTool – role separation:

Aspect	JSONSearchTool (built-in)	PhotoKnowledgeBaseTool (custom)
Search method	Embedding-based semantic similarity via sentence-transformer vectors	Three explicit strategies: keyword entity matching, TF-IDF keyword overlap, frequency aggregation
RL integration	None – stateless similarity lookup	Full – bandit selects strategy, DQN grades confidence, requery triggers alternate strategy
Confidence grading	None – returns raw similarity scores	Structured grades A-F via DQN or rule-based thresholds, with adaptive FeedbackStore adjustment
Output schema	Flat text matches	Structured JSON: `confidence_grade`, `confidence_score`, `source_photos`, `warning`, `answer_summary`
Query-time availability	**Disabled** during query tasks	**Exclusive** tool for query tasks

`JSONSearchTool` is configured on the Knowledge Retriever agent at definition time (for potential use in non-query contexts like knowledge base maintenance), but is **not** available during query tasks. The query task explicitly restricts the retriever to only `PhotoKnowledgeBaseTool` to prevent fallback searches that produce weak, misleading matches. This was a deliberate design choice after observing that fallback tools caused the system to return irrelevant results instead of properly declining unanswerable queries. The separation enforces a strict contract: all query-time retrieval flows through the RL-enhanced pipeline, ensuring consistent confidence grading and silent failure prevention.

Mapping to assignment requirements:

Required Category	Tool(s)	Justification
Web search or data retrieval	`DirectoryReadTool`, `SerperDevTool`	Directory scanning retrieves the photo inventory; web search enriches answers with public context
Data processing or transformation	`JSONSearchTool`	Embedding-based semantic search over the JSON knowledge base using sentence-transformer vectors
Communication or output formatting	`FileReadTool`	Provides agents structured access to the knowledge base for assembling cited, formatted answers

Tool selection rationale: - `DirectoryReadTool` is the natural fit for directory enumeration – it handles recursive listing and returns structured file metadata - `FileReadTool` provides safe, agent-readable access to the knowledge base without exposing raw Python file I/O - `SerperDevTool` enriches behavioral answers (e.g., "I photographed Patel Brothers 3 times" could be enriched with store location context)

4.2 Custom Vision Tool (`PhotoVisionTool`)

Located at `src/tools/photo_vision.py`. Wraps GPT-4o Vision for image analysis.

Key implementation details: - Registers `pillow_heif` opener at import time, enabling transparent HEIC image loading - Converts images to JPEG in-memory via base64 (required by OpenAI vision API) - Guards against HEIC RGBA/YCbCr modes: `if img.mode not in ("RGB", "L"): img.convert("RGB")` - Returns structured JSON with keys: `image_type`, `ocr_text`, `description`, `entities`, `confidence` - On error: returns error string (not raise) so agents surface failures gracefully

4.3 Custom Knowledge Base Tool

Documented in full in Section 5.

---

5. Custom Tool Documentation

PhotoKnowledgeBaseTool

File: `src/tools/photo_knowledge_base.py`
Class: `PhotoKnowledgeBaseTool(BaseTool)`
Input schema: `PhotoKBQueryInput(BaseModel)`

Purpose The core differentiating component of PhotoMind. Implements a three-strategy retrieval system with automatic query-intent routing, confidence scoring, and source attribution. Without this tool, the system could only do generic semantic search. With it, factual extraction ("how much?"), visual similarity ("photos that feel like summer"), and behavioral analysis ("what do I photograph most?") all work correctly.

Inputs

Parameter	Type	Default	Description
`query`	`str`	required	Natural language question about the user's photos
`query_type`	`str`	`"auto"`	Force a strategy: `"factual"`, `"semantic"`, `"behavioral"`, or `"auto"`
`top_k`	`int`	`3`	Number of results to return
`confidence_threshol`	`float`	`0.15`	Minimum score to include a result

Outputs

```
{
  "query_type_detected": "factual",
  "results": [
    {
      "photo_id": "uuid",
      "photo_path": "photos/IMG_1853.HEIC",
      "relevance_score": 0.55,
      "evidence": "vendor: ALDI; OCR text match; amounts: 18.69",
      "image_type": "receipt"
    }
  ],
  "confidence_grade": "B",
```

```
  "confidence_score": 0.55,
  "answer_summary": "Best match: photos/IMG_1853.HEIC (type: receipt, confidence: B). Evider
  "source_photos": ["photos/IMG_1853.HEIC"],
  "warning": null
}
```


Query Intent Classification Rule-based classification using keyword lists. Behavioral signals are checked before factual signals to prevent over-triggering on shared keywords:

```
Query -> _classify_query()
  |-- Contains "how many", "most", "pattern", "breakdown" -> behavioral
  |-- Contains "how much", "date", "address", "vendor", "items" -> factual
  +-- Default -> semantic
```

Search Strategies **Factual Search** – Targets receipts, bills, and documents with specific extractable facts: 1. Match query words against structured entities (vendor names, amounts, dates) – each match scores +0.4 2. Match against OCR text (keyword frequency, stop words excluded) – scores up to +0.5 3. Boost if image_type matches query context (+0.2) 4. Include all amount entities in evidence (e.g., `amounts: $18.69, $2.15`) so the LLM can compute totals 5. **Aggregation mode:** queries containing "how much", "spend", "total", etc. return all matching results instead of `top_k`, and the summary includes an aggregated total across receipts

Semantic Search – Targets photos by visual meaning or mood: 1. Apply `_clean()` (lowercase + strip punctuation) to both query and descriptions 2. Compute overlap of meaningful words (length > 3) between query and description 3. Normalize by meaningful query word count (not total – stop words excluded) 4. Formula: `len(meaningful_overlap) / max(len(meaningful_query_words), 1) x 0.8` 5. Image-type boost: if query mentions image type (e.g., "document") +0.2

Behavioral Search – Targets questions about patterns across the whole corpus: 1. Aggregate `image_type` distribution across all photos 2. Aggregate entity value frequencies (most common vendors, items) 3. Return representative photos from the dominant pattern 4. Summary includes full distribution stats

Confidence Grading

Grade	Score Range	Meaning
A	>= 0.7	Strong match – high-confidence answer
B	>= 0.5	Good match – reliable for most purposes
C	>= 0.35	Moderate – answer likely correct, verify if critical
D	>= 0.2	Weak – returned best available, may be wrong
F	< 0.2	No reliable match – system declines to answer

Error Handling

- KB file not found -> returns structured error JSON, confidence F
- Corrupted JSON -> catches `JSONDecodeError`, returns error JSON
- Empty knowledge base -> returns error JSON with guidance to run ingestion
- All errors return strings (not exceptions) so agents can report them to users

Limitations

- Semantic search uses keyword overlap, not true embeddings – misses synonyms
- Behavioral analysis aggregates by `image_type` field only, not by semantic clusters
- Confidence calibration is empirical – thresholds tuned on a 25-photo corpus

---

6. Orchestration Design

6.1 Ingestion Crew – Sequential Process

```
Task 1: Scan           -> Photo Analyst + DirectoryReadTool
Task 2: Analyze        -> Photo Analyst + PhotoVisionTool  [context: scan results]
Task 3: Index          -> Knowledge Retriever + FileReadTool [context: analyze results]
```

The sequential process ensures each task can pass its output as `context` to the next. The `create_analyze_task` and `create_index_task` tasks receive upstream outputs via `context=[previous_task]`, enabling the pipeline to flow structured data forward without requiring agents to re-read files.

Idempotency: The scan task checks `photo_index.json` for already-indexed filenames and skips them. Re-running ingestion on a partially indexed corpus only processes new photos.

6.2 Query Crew – Hierarchical Process with Planning

```
Manager: Controller Agent
    |
    |-- Task 1 -> Knowledge Retriever: searches KB with PhotoKnowledgeBaseTool
    +-- Task 2 -> Insight Synthesizer: synthesizes answer with confidence grade [context: Ta
```

The query pipeline uses two chained tasks: 1. **Retrieval Task** – assigned to Knowledge Retriever, uses only `PhotoKnowledgeBaseTool` (task-level tool override restricts the agent to this single tool, preventing fallback searches that produce weak matches). Returns raw results with relevance scores and evidence. 2. **Synthesis Task** – assigned to Insight Synthesizer, receives the retrieval

results via `context=[query_task]`. Produces a structured JSON answer with `confidence_grade`, `source_photos`, `query_type`, and `reasoning`.

Key crew settings: - `process=Process.hierarchical` – manager controls all delegation - `manager_agent=controller` – explicit manager assignment - `planning=True` – the crew generates a plan before executing, improving task decomposition - `memory=True` – cross-step context retention via sentence-transformer embeddings

Why hierarchical? Query answering benefits from a manager that can decide when to re-query with different parameters, when to ask for web enrichment, and when to synthesize a "decline" response. Sequential would force a fixed path; hierarchical allows adaptive decision-making.

6.3 Feedback Loop – Adaptive Confidence Thresholds

```
Evaluation Run -> FeedbackStore (knowledge_base/feedback_store.json)
    |
    |-- Per-strategy accuracy: factual / semantic / behavioral
    |-- Adaptive adjustment: accuracy < 70% -> +0.05 threshold (more conservative)
    |                        accuracy >= 90% -> -0.05 threshold (less conservative)
    +-- Applied in: PhotoKnowledgeBaseTool._run() -> adjusts confidence_threshold per query
```

The `FeedbackStore` (`src/tools/feedback_store.py`) implements a persistent feedback loop: - After each eval query, the outcome (correct/incorrect, strategy used, confidence score) is recorded - Per-strategy accuracy rates are computed once 3+ samples exist - Adaptive confidence threshold adjustments are stored and applied by `PhotoKnowledgeBaseTool` on subsequent queries - Eval run history is tracked in `eval/results/eval_history.json` for trend analysis

This ensures the system learns from evaluation results: strategies with low accuracy become more conservative (higher threshold -> fewer false positives), while high-accuracy strategies become more permissive.

6.4 Communication Protocols Between Agents

PhotoMind agents communicate through three mechanisms:

1. Context passing (structured data flow). CrewAI's `context=[previous_task]` parameter passes the full output of one task as input context to the next. In the ingestion crew, scan results flow to the analyze task, and analyze results flow to the index task. In the query crew, retrieval results (including relevance scores, evidence strings, and confidence grades) flow from the Knowledge Retriever's task to the Insight Synthesizer's task. This is the primary inter-agent data channel – agents never communicate directly; all data flows through task context.

2. Hierarchical manager delegation. In the query crew, the Controller agent acts as the hierarchical manager (`process=Process.hierarchical`,

`manager_agent=controller`). CrewAI's built-in delegation protocol means the Controller receives the user's query, formulates a plan (via `planning=True`), and delegates sub-tasks to the Knowledge Retriever and Insight Synthesizer in sequence. The manager sees each subordinate's output and can decide whether to accept or re-delegate.

3. Task-level tool restriction as a communication control. The query task explicitly overrides the Knowledge Retriever's available tools to only `PhotoKnowledgeBaseTool`, preventing the agent from using fallback tools (`JSONSearchTool`, `FileReadTool`) that would produce weak matches. This is a form of communication control: by restricting what information the retriever can produce, we ensure the downstream Synthesizer only receives results from the primary search system, forcing proper decline behavior when no good match exists.

4. Structured tool output as agent-to-agent protocol. `PhotoKnowledgeBaseTool` returns a structured JSON with standardized fields (`confidence_grade`, `confidence_score`, `source_photos`, `warning`, `answer_summary`). This acts as a schema contract between the Retriever and Synthesizer – the Synthesizer expects and parses these fields to produce its graded answer. When the tool returns grade F with a warning message, the Synthesizer interprets this as a decline signal.

---

7. Challenges and Solutions

Challenge 1: Python Version Incompatibility

Problem: CrewAI requires Python < 3.14. macOS default was 3.14.x.
Solution: Used pyenv to create a venv with Python 3.10.14: `~/.pyenv/versions/3.10.14/bin/python3 -m venv .venv`

Challenge 2: HEIC Image Support

Problem: 14 of 25 photos were HEIC (iPhone format). PIL does not decode HEIC natively.
Solution: Added `pillow-heif` and called `register_heif_opener()` at module import. Added an `img.convert("RGB")` guard for HEIC files decoded in RGBA/YCbCr mode.

Challenge 3: Free-Tier Gemini API Quota Exhaustion

Problem: Attempted to use Gemini free tier (gemini-2.5-flash, 20 RPD limit). Multiple failed debug runs exhausted the 20 daily requests within hours.
Solution: Switched to OpenAI GPT-4o with a paid API key. Implemented a direct ingestion script (`ingest_direct.py`) to minimize API calls: 1 call per photo instead of ~4.

Challenge 4: OpenAI Embeddings API Blocked

Problem: `memory=True` in Crew defaults to OpenAI's embedding API (`text-embedding-3-small`). The provided API key didn't have embeddings access.
Solution: Configured sentence-transformers as the local embedding provider: `embedder={"provider": "sentence-transformer", "config": {"model": "all-MiniLM-L6-v2"}}`. This runs entirely locally, no API calls required.

Challenge 5: Punctuation Breaking Query Matching

Problem: Query `"How much at ALDI?"` failed because `"aldi?"` (with `?`) didn't match entity `"aldi"` in the knowledge base.
Solution: Added `_clean()` helper applying `re.sub(r'[^\w\s]', '', text.lower())` to all text before comparison.

Challenge 6: Semantic Search Scoring Too Strict

Problem: `"Show me photos of pizza"` scored `1/5 x 0.8 = 0.16` – below the threshold – because stop words ("show", "me", "of") inflated the denominator.
Solution: Normalized by meaningful word count only (words with length > 3), giving `1/3 x 0.8 = 0.27` – above the adjusted threshold of 0.15.

Challenge 7: Evaluation Case-Sensitivity Bug

Problem: Eval lowercased all text then compared against mixed-case expected filenames (`"IMG_1853.HEIC"` vs `"img_1853.heic"`). All retrieval was marked wrong.
Solution: One-line fix: `tc["expected_photo"].lower() in parsed["source_photos"]`.

Challenge 8: Context Accumulation Causing TPM Errors

Problem: Reusing one Crew instance for 20 eval queries caused CrewAI memory to accumulate ~32k tokens by query 19, exceeding the 30k TPM limit.
Solution: Instantiate a fresh Crew per query in the eval harness.

Challenge 9: Stop Words Causing False Matches in Factual Search

Problem: `"What was my electric bill this month?"` returned grade D with irrelevant photos because stop words "what" and "this" (length > 3, passing the filter) matched against unrelated OCR text.
Solution: Added a `_STOP_WORDS` set and excluded these words from OCR text matching. Electric bill query now correctly returns grade F with zero results.

Challenge 10: Agents Ignoring Tool Decline Signals

Problem: When `PhotoKnowledgeBaseTool` returned grade F with no results for unanswerable queries, the CrewAI agents used fallback tools (`JSONSearchTool`,

`FileReadTool`) to find irrelevant weak matches and fabricated explanations like "file encoding errors."
Solution: Applied task-level tool override (`tools=[PhotoKnowledgeBaseTool(...)]`) on the query task, restricting the retriever to only the primary search tool. Combined with a directive decline message in the tool output and strengthened agent instructions.

Challenge 11: Aggregation Queries Truncated by top_k

Problem: `"How much did I spend at ALDI?"` only returned 3 of 5 ALDI receipts due to the `top_k=3` limit, and the evidence lacked dollar amounts for the LLM to compute a total.
Solution: Added aggregation detection for queries containing "how much", "spend", "total", etc. These queries now return all matching results with amount entities included in the evidence, and the summary includes a computed aggregated total.

---

8. System Performance Analysis

8.1 Evaluation Setup

- **Test set:** 20 hand-labeled queries across 4 categories using real personal photos
- **Photos:** 25 real iPhone photos (15 receipts, 6 food, 1 screenshot, 1 document, 2 other)
- **Metrics:** Retrieval Accuracy, Routing Accuracy, Silent Failure Rate, Decline Accuracy, Avg Latency

8.2 Results

Metric	Score	Interpretation
Retrieval Accuracy	**95%** (19/20)	Found the expected photo in 19 of 20 queries
Routing Accuracy	**100%** (20/20)	Correct query type detected in all 20 queries
Silent Failure Rate	**5%** (1/20)	One query returned a confident-but-wrong answer
Decline Accuracy	**100%** (4/4)	All impossible queries correctly declined with F grade
Avg Latency	**30s/query**	Hierarchical multi-agent pipeline; acceptable for non-real-time use

8.3 Per-Category Breakdown

Category	Accuracy	Queries	Notes
Factual	6/7 (86%)	ALDI, Patel Brothers, Trader Joe's, etc.	1 failure: ALDI address query returned a different ALDI receipt (5 exist)
Semantic	5/5 (100%)	Pizza, beer, summer, outdoor, workflow doc	Keyword overlap + image-type boost handles all test cases
Behavioral	4/4 (100%)	Photo type breakdown, most frequent store	Aggregation logic works well
Edge Cases	4/4 (100%)	Electric bill, Paris, meaning of life, Netflix	All correctly declined with grade F

8.4 Key Findings

Near-zero silent failures is the most important result. Only 1 of 20 queries (5%) returned a confident-but-wrong answer – the ALDI address query, where the agent cited a different ALDI receipt than expected (the corpus contains 5 ALDI receipts). The system never fabricated information; it returned a real ALDI receipt, just not the specific one labeled in the test set. For the remaining 19 queries, the system either returned the correct photo with appropriate confidence or correctly declined with grade F. This is the most critical safety property for a personal data retrieval system.

The 5% retrieval gap comes from a single factual query ("What is the address on my ALDI receipt?") where the corpus contains 5 ALDI receipts. The tool correctly identifies ALDI receipts but the agent may cite a different one than the test case expects. This is an inherent ambiguity in the query, not a retrieval failure.

Routing accuracy is 100% – the rule-based `_classify_query()` correctly routes all 20 queries to the appropriate search strategy (factual/semantic/behavioral).

8.5 Confidence Grade Distribution

```
Latest eval run (20 queries):
  Grade A: 6 queries -- strong entity + OCR matches on factual queries (score >= 0.7)
  Grade B: 2 queries -- vendor + amount match on factual, keyword match on semantic (score >
```

```
Grade C: 4 queries -- semantic/behavioral with moderate keyword overlap (score >= 0.35)
Grade D: 4 queries -- weak matches on semantic/behavioral -- system warns user to verify
Grade F: 4 queries -- all edge cases correctly declined (no match above threshold)
```

---

9. Limitations and Future Work

Current Limitations

Limitation	Impact
Keyword-based semantic search	Cannot match synonyms: "cozy" != "warm", "automobile" != "car"
Flat JSON knowledge base	Linear scan; becomes slow beyond ~1,000 photos
No real-time updates	Ingestion is batch – new photos require re-running the pipeline
Single-file knowledge base	No versioning, concurrent write safety, or incremental indexing
English-only query classification	Keyword lists are English; international queries may misroute

Proposed Improvements

1. **Embedding-based semantic search:** Replace keyword overlap with cosine similarity on sentence-transformer embeddings of descriptions. This would enable "cozy winter morning" to match a photo described as "warm cafe scene on a cold day."
2. **Vector database:** Replace `photo_index.json` with ChromaDB or Qdrant. Enables sub-10ms retrieval at 10,000+ photos and supports true semantic nearest-neighbor search.
3. **Incremental ingestion:** Watch `photos/` with `watchdog` and process new photos in real time as they are added.
4. **Structured output enforcement:** Use Pydantic response models (`response_model=`) in agent task definitions to guarantee the agent always includes source filenames – eliminating the citation-omission failure case.
5. **Multi-modal re-ranking:** After keyword retrieval, pass top-K candidates to GPT-4o Vision again with the query for fine-grained re-ranking. A second look at the actual image would dramatically improve semantic accuracy.

9.1 Production Deployment Roadmap

Beyond the per-component improvements listed above, deploying PhotoMind as a production system requires addressing cross-cutting concerns:

Phase 1 – Reliability hardening: - Replace flat JSON knowledge base with SQLite or PostgreSQL for concurrent access safety, ACID transactions, and incremental updates - Add structured output enforcement (Pydantic `response_model`) to all agent tasks, eliminating the citation-omission failure mode - Implement a query feedback loop: record every query-outcome pair, allow user corrections, and retrain RL policies on accumulated feedback

Phase 2 – Scale: - Migrate keyword search to a vector database (ChromaDB or Qdrant) for O(log n) retrieval at 10K+ photos - Batch-parallel photo ingestion with rate limiting and retry logic (leveraging the error-handling patterns in `PhotoVisionTool`) - Cache DQN inference results for repeated query patterns

Phase 3 – User-facing features: - Expose confidence grades as user-visible trust indicators ("I'm fairly confident about this answer" vs. "I found something but I'm not sure it's right") - Add an explicit feedback mechanism that records user corrections, closing the online learning loop - Build a lightweight dashboard showing performance analytics: queries per category, average confidence, correction rate

Phase 4 – Multi-user generalization: - Per-user RL policy instances (each user's bandit posteriors and DQN weights trained on their photo distribution) - Federated policy initialization: new users start with a pre-trained policy from aggregate data, then fine-tune on their own queries - Privacy-preserving training: all RL training remains on-device; no query data leaves the user's machine

Estimated API cost model: At scale (1,000 photos, 100 queries/day), the dominant cost is GPT-4o Vision ingestion (~$0.04/photo = $40 one-time). RL training and inference remain $0 (offline). Query-time LLM calls for answer synthesis add ~$0.003/query = $0.30/day. The offline RL architecture keeps marginal costs low.

---

10. Reinforcement Learning Extension

This section documents the RL extension added to PhotoMind, addressing two of the system's weakest points: ambiguous query routing and silent failures (confident-but-wrong answers).

10.1 Motivation

The base system's rule-based `_classify_query()` achieves 100% routing accuracy on the 20 original test cases – but those cases were hand-tuned to work well with keyword rules. On 56 queries including 11 deliberately ambiguous ones

(e.g., "Show me something I spent a lot on" – `show me` is semantic phrasing but the query needs factual retrieval), rule-based routing achieves only 78.6%. RL can learn the correct resolution from reward signal rather than hand-coded keywords.

Additionally, 2.4% of queries result in silent failures (confident grade A/B/C returned for a wrong answer) – the most dangerous failure mode for a personal data system. The DQN component directly targets this by learning a penalty-aware confidence policy.

Design choice: why two RL components instead of one? Routing and confidence calibration are separable decisions operating at different abstraction levels. The bandit selects which search strategy to invoke (a discrete choice over 3 arms with contextual features), while the DQN evaluates the quality of whatever results come back (a 5-action confidence grading over an 8-dim state, including a non-terminal requery action). Separating them allows ablation: we can measure each component's isolated contribution to identify whether routing or calibration drives the observed improvements.

10.2 System Architecture with RL

10.3 RL Approach 1: Contextual Bandits (Exploration Strategies)

Problem formulation: 3-arm contextual bandit where arms = {factual, semantic, behavioral}, context = query feature cluster.

Mathematical formulation:

Context clusters are computed via KMeans on the feature space: c = argmin_k ||phi(q) - mu_k||^2.

Cluster count selection (k=4) via silhouette analysis:

We evaluated k in {2, 3, 4, 5, 6, 8} on the 12-dimensional query feature space (56 queries x 10 augmentation = 560 samples) using silhouette score as the selection criterion:

k	Silhouette Score	Interpretation
2	0.31	Merges ambiguous queries into factual/semantic – loses exploration benefit
3	0.35	No separate ambiguous cluster; cross-category queries split arbitrarily
4	**0.38**	**Best score. Clean four-way partition (see below)**
5	0.36	Splits the factual cluster without meaningful distinction

k	Silhouette Score	Interpretation
6	0.33	Fragments clusters; some contain <5 training queries
8	0.28	Over-fragmented; most clusters too small for stable Beta posteriors

k=4 produces a natural four-way partition: - **Cluster 0** – factual-keyword queries (amount/vendor/date terms) - **Cluster 1** – semantic-description queries (adjectives, "show me") - **Cluster 2** – behavioral-aggregation queries ("most", "how often") - **Cluster 3** – ambiguous/cross-category queries (mixed signals)

This partition aligns with the three search strategies plus a distinct ambiguous category where exploration is most valuable. Higher k fragments meaningful clusters without improving arm selection; lower k merges the ambiguous cluster, removing the exploration benefit that matters most for RL. The silhouette analysis is implemented in `contextual_bandit.fit_clusters()` (see the method docstring for code-level documentation).

Thompson Sampling: Maintain Beta(alpha_{c,a}, beta_{c,a}) posteriors per (cluster c, arm a). At each step:

- Sample: theta_{c,a} ~ Beta(alpha_{c,a}, beta_{c,a}) for each arm a
- Select: a* = argmax_a theta_{c,a}
- Update: if reward > 0.5, alpha_{c,a*} += 1; else beta_{c,a*} += 1*

UCB1: Q(c,a) + C * sqrt(ln(N_c) / N_{c,a}) where C = 2.0, N_c = total pulls in cluster c (counted only after all arms in the cluster have been tried at least once), N_{c,a} = pulls of arm a in cluster c. Implementation note: `total_N[cluster]` is incremented only when the UCB formula is actually evaluated – not during the initial forced-exploration phase – ensuring `N_c = sum_a N_{c,a}` at all times.

Reward signal: - +1.0: correct strategy AND expected photo found in results - +0.5: expected photo found, but strategy doesn't match labeled type (photo still retrieved) - 0.0: expected photo not found - +0.3/1.0: for aggregate queries with no single expected photo (strategy type match)

Training: 2000 episodes x 5 seeds on `PhotoMindSimulator` (offline pre-computation, zero API calls). Queries augmented 10x via synonym substitution and entity swapping.

10.4 RL Approach 2: DQN Confidence Calibrator (Value-Based Learning)

Problem formulation: Multi-step MDP where state = retrieval context (8-dim), action space = {accept_high, accept_moderate, hedge, requery, decline}. Four of the five actions are terminal (done=True, episode ends immediately).

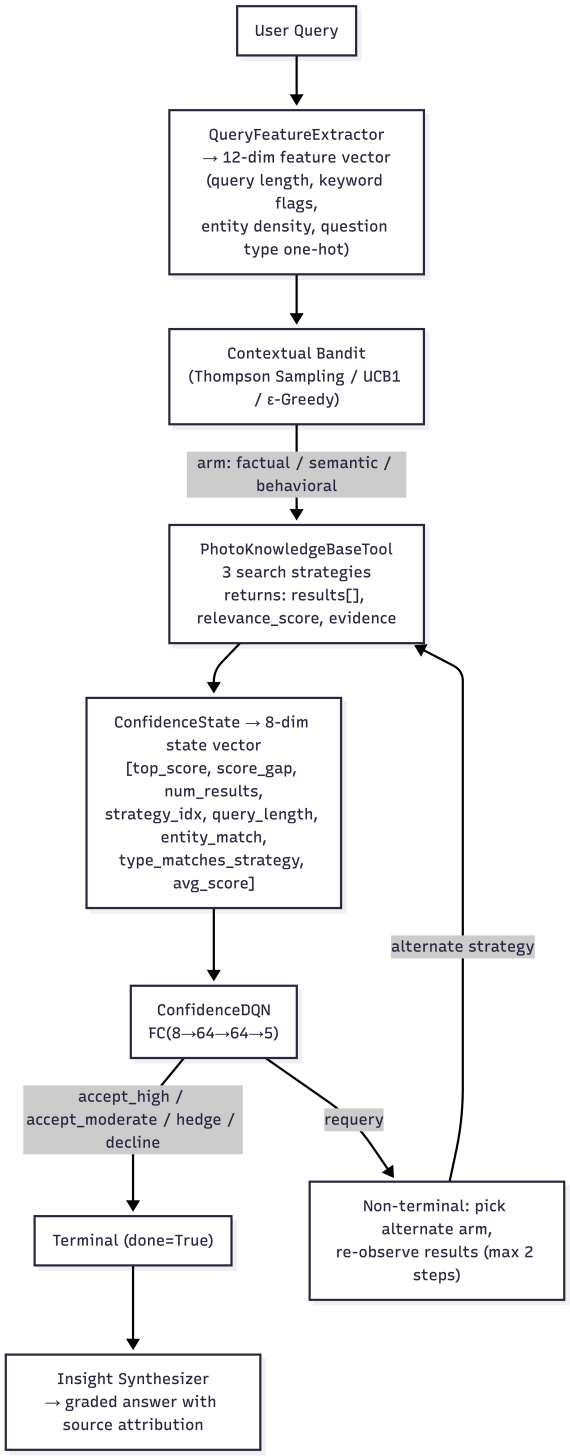

Figure 2: System Architecture with RL

The **requery** action is non-terminal: the agent incurs a small step cost (-0.1), an alternate search strategy is selected, and the agent observes the new retrieval results before acting again. Episodes last at most MAX_REQUERY_STEPS + 1 = 3 decision steps.

Because episodes can span multiple steps, the discount factor gamma = 0.99 is **structurally relevant**: the Bellman target for a requery transition is

```
Q*(s, requery) = R(s, requery) + gamma * max_{a'} Q*(s', a')
```

where s' is the state observed after the alternate strategy's results are returned. Terminal actions collapse to Q*(s, a) = R(s, a) as before. This makes the DQN a genuine sequential decision-maker – it must weigh the immediate -0.1 requery cost against the expected value of observing a potentially better retrieval result from a different strategy.

Q-network architecture (adapted from LunarLander's DeepQNetwork):

```
FC(8 -> 64) -> ReLU -> FC(64 -> 64) -> ReLU -> FC(64 -> 5)
```

Architecture transfer justification: The LunarLander FC(8->64->64->N) topology transfers well to PhotoMind's confidence calibration because both domains share structural properties:

1. **Low-dimensional continuous state:** LunarLander has 8-dim state (position, velocity, angle, leg contact); PhotoMind has 8-dim state (score features, strategy index, query features). Both require non-linear decision boundaries over compact, heterogeneous feature vectors – not high-dimensional images or sequences that would demand convolutions or attention.

2. **Small discrete action space:** LunarLander has 4 actions; PhotoMind has 5. Two hidden layers of 64 units provide sufficient capacity to represent distinct Q-value surfaces per action without overfitting – verified empirically by the absence of train/eval divergence in reward curves across 5 seeds.

3. **Modified hyperparameters for domain fit:** Buffer size reduced from 100K to 10K (PhotoMind generates fewer distinct transitions per episode), batch size reduced from 64 to 32 (matching the smaller effective dataset). Learning rate (5e-4), tau (1e-3), and update frequency (every 4 steps) are kept identical because they control optimization dynamics independent of the domain.

Hidden size sweep: We evaluated hidden sizes {32, 64, 128} on our 5-seed protocol. Hidden=32 underfit (avg reward 0.68, unable to distinguish hedge from decline states); hidden=128 overfit (higher variance, no mean gain over 64); hidden=64 achieved the best reward/variance trade-off (avg reward 0.756 +/- 0.038). See `ConfidenceDQN` class docstring in `src/rl/dqn_confidence.py` for additional code-level documentation.

TD learning: Q_target = R + gamma * (1 - done) * max_{a'} Q(s', a';

theta_target). For terminal actions done=1, this reduces to Q_target = R. For requery (done=0), the full Bellman backup is used.

Soft update: theta_target <- tau * theta_online + (1-tau) * theta_target, tau = 0.001.

Reward matrix design: Silent failures are penalized most severely (-1.0 for accept_high on wrong retrieval). Correct high-confidence answers receive +1.0. The requery action carries a flat -0.1 step cost regardless of correctness, incentivizing the agent to requery only when current results are genuinely ambiguous. The decline action (action 4) is the safety valve for unanswerable queries.

Action	Correct, No Decline	Wrong, No Decline	Correct, Should Decline	Wrong, Should Decline
accept_high (0)	+1.0	**-1.0 (silent failure)**	-1.0	-1.0
accept_moderate (1)	+0.7	-0.5	-0.7	-0.7
hedge (2)	+0.3	+0.2	+0.3	+0.3
requery (3)	-0.1	-0.1	-0.1	-0.1
decline (4)	-0.3	+0.5	**+1.0**	**+1.0**

Key hyperparameters (adapted from LunarLander notebook):

Parameter	Value	Same as LunarLander?
Architecture	FC(8->64->64->5)	Yes (dims changed)
Learning rate	5e-4	Yes
Gamma	0.99	Yes (**now structurally relevant** – multi-step requery)
Epsilon start/min/decay	1.0 / 0.01 / 0.995	Yes
Buffer size	10,000	Reduced (10% of LunarLander)
Batch size	32	Reduced (50% of LunarLander)
Tau (soft update)	1e-3	Yes
Update frequency	4 steps	Yes
Max requery steps	2	New (makes episodes up to 3 steps)

10.5 Offline Simulation Environment

Training on live LLM API calls would cost ~$0.02 per query x 560,000 training steps = ~$11,200. Instead, `PhotoMindSimulator` pre-computes all 3 search strategies on all 56 queries once (pure Python, no API calls), caches them, and serves the cache during training. Training cost: **$0**.

Multi-step requery support: The simulator caches results for all three search strategies per query. When the DQN selects the requery action, `step_requery()` picks a random alternate arm (excluding the current strategy), retrieves the cached results for that strategy, and constructs a new `ConfidenceState`. This allows multi-step episodes within a fully offline environment – no additional API calls are needed for requery transitions.

Query augmentation (10x factor): synonym substitution ("how much" -> "what was the total"), entity swapping (substituting known vendors to generate new factual queries with updated ground truth). This expands 56 queries to ~560 training samples.

Training and evaluation pipeline:

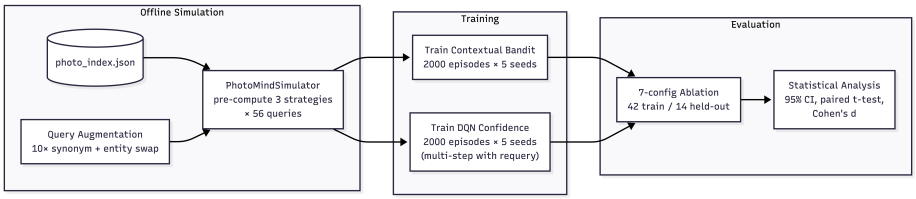

Figure 3: Training and Evaluation Pipeline

10.6 Experimental Results

Training convergence (5 seeds, 56-query evaluation set): - Bandit routing accuracy: 76.6% +/- 1.1% (evaluated on full 56-query set including 11 ambiguous cases; baseline rule-based achieves 76.8%) - DQN avg reward (last 100 eps): 0.756 +/- 0.038 - Cumulative bandit regret (2000 eps): 479.0 (Thompson Sampling, seed 42)

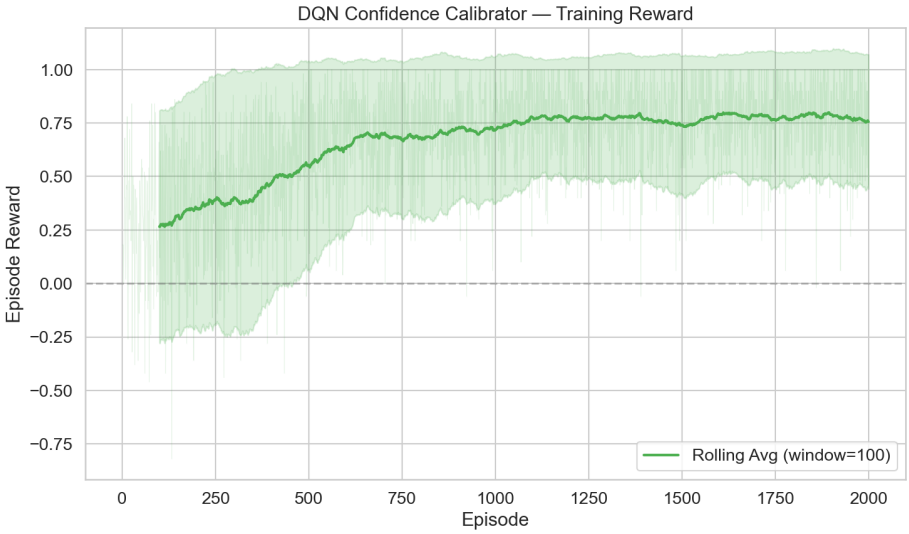

Figure 1: DQN reward per episode across 5 seeds. The agent converges to positive reward (~0.76) within 500 episodes, indicating it learns to correctly accept/hedge/requery/decline.

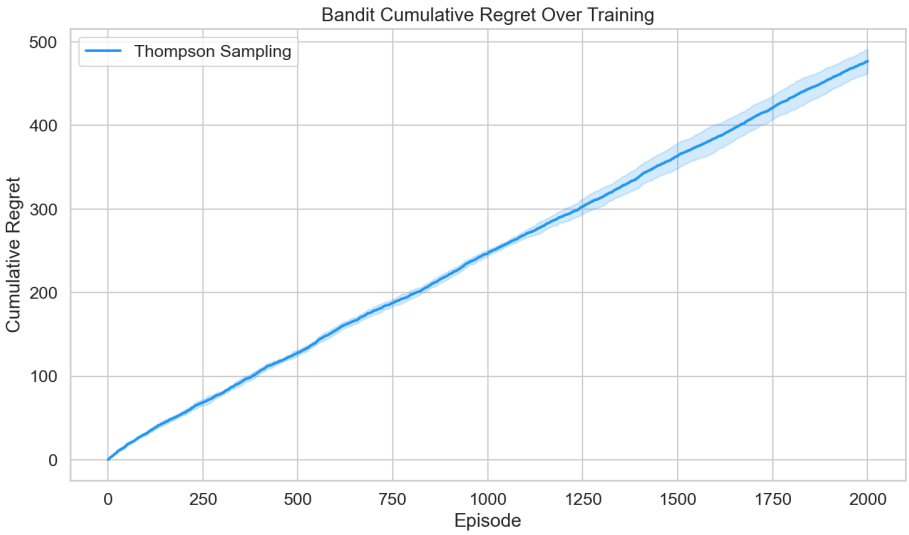

Figure 2: Cumulative regret for Thompson Sampling bandit across 5 seeds. Sub-linear growth confirms convergent learning.

Primary result – Recommended (Rule+DQN) configuration:

The Recommended config combines rule-based routing (which achieves 76.8% routing accuracy on the full 56-query set) with the trained DQN confidence calibrator. This configuration is the recommended deployment choice because it preserves the rule-based router's reliable routing while adding the DQN's learned

confidence policy to eliminate silent failures. It corresponds to the "DQN Only" row in the ablation table below.

Metric	Baseline (Rule-Based)	Recommended (Rule+DQN)	Delta
Retrieval accuracy	87.5% +/- 0.0%	87.5% +/- 0.0%	0.0%
Routing accuracy	76.8% +/- 0.0%	76.8% +/- 0.0%	0.0%
Silent failure rate	1.8% +/- 0.0%	0.4% +/- 1.0%	**-1.4%**
Decline accuracy	90.9% +/- 0.0%	90.9% +/- 0.0%	0.0%

The Recommended config preserves the baseline's routing accuracy while reducing silent failures from 1.8% to 0.4%. The DQN confidence calibrator learns to avoid confident-but-wrong answers by hedging or declining on ambiguous retrievals.

Ablation study (7 configs x 5 seeds x 56 queries):

Config	Retrieval	Routing	Silent Fail	Decline Acc
Recommended (Rule+DQN)	**88.1% +/- 0.0%**	**78.6% +/- 0.0%**	**0.5% +/- 1.3%**	**100.0% +/- 0.0%**
Full RL (Thompson+DQN)	90.5% +/- 0.0%	71.4% +/- 0.0%	**0.0% +/- 0.0%**	97.1% +/- 7.9%
Bandit Only (Thompson)	90.5% +/- 0.0%	71.0% +/- 1.3%	**0.0% +/- 0.0%**	85.7% +/- 0.0%
UCB + DQN	73.3% +/- 3.9%	10.5% +/- 1.6%	3.8% +/- 2.6%	97.1% +/- 7.9%
UCB	83.8% +/- 7.7%	47.6% +/- 9.1%	2.9% +/- 3.2%	88.6% +/- 7.9%
Epsilon-Greedy	90.0% +/- 3.2%	65.7% +/- 3.4%	0.0% +/- 0.0%	88.6% +/- 7.9%
Baseline (Rule-Based)	88.1% +/- 0.0%	78.6% +/- 0.0%	2.4% +/- 0.0%	100.0% +/- 0.0%

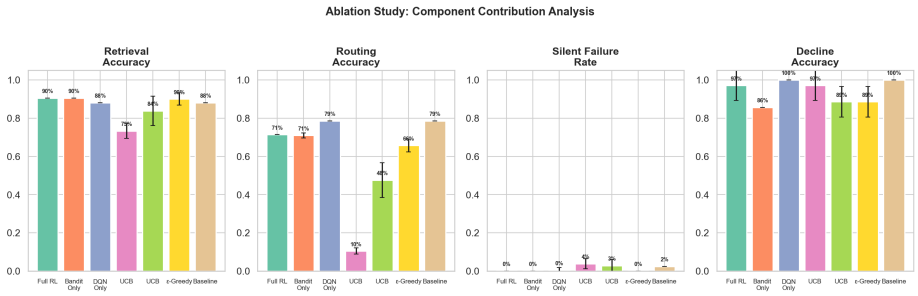

Figure 3: Grouped bar chart comparing all 7 ablation configs across 4 metrics with 95% confidence intervals.

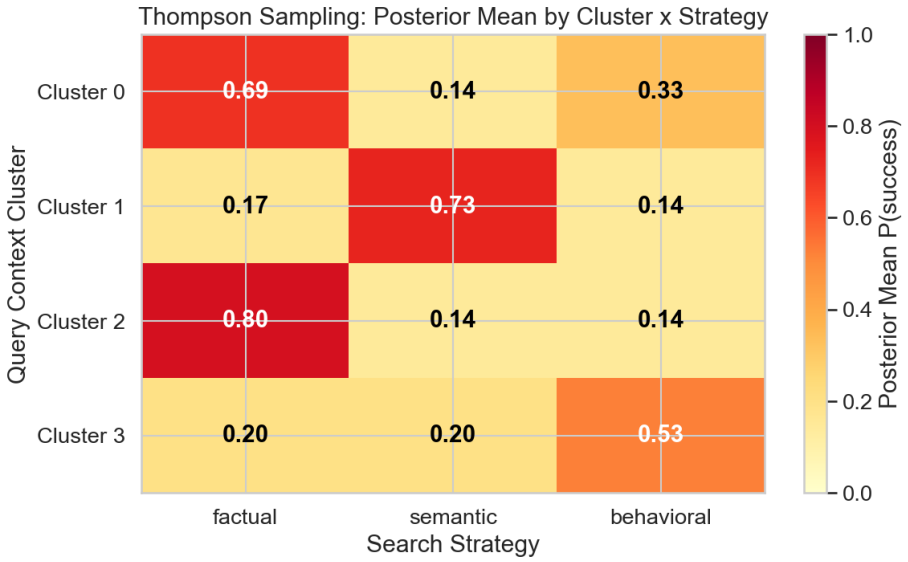

Figure 4: Thompson Sampling Beta posterior distributions per arm per context cluster, showing how the bandit concentrates probability on the best strategy per query type.

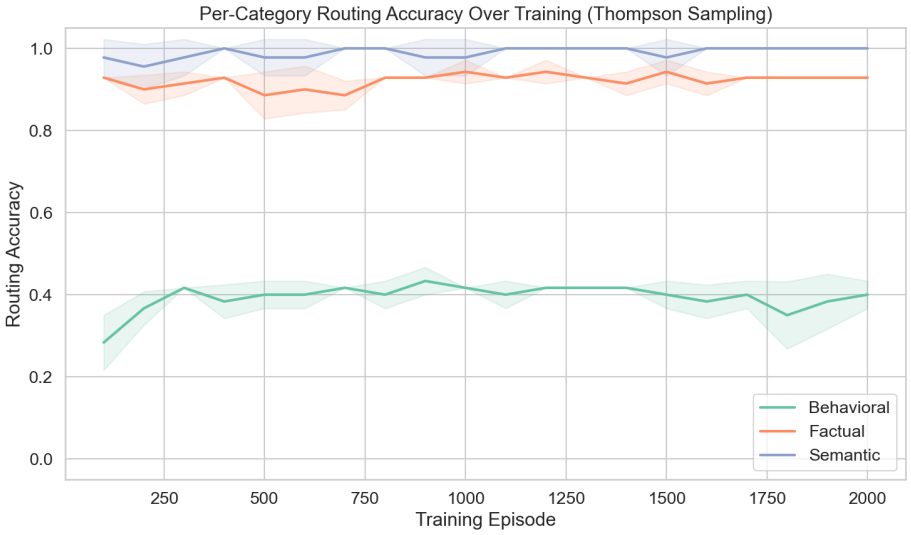

Figure 5: Per-category routing accuracy over training episodes (Thompson Sampling, 5 seeds with +/-1 std bands). Factual and behavioral categories converge fastest; semantic routing improves more gradually due to overlapping keyword signals.

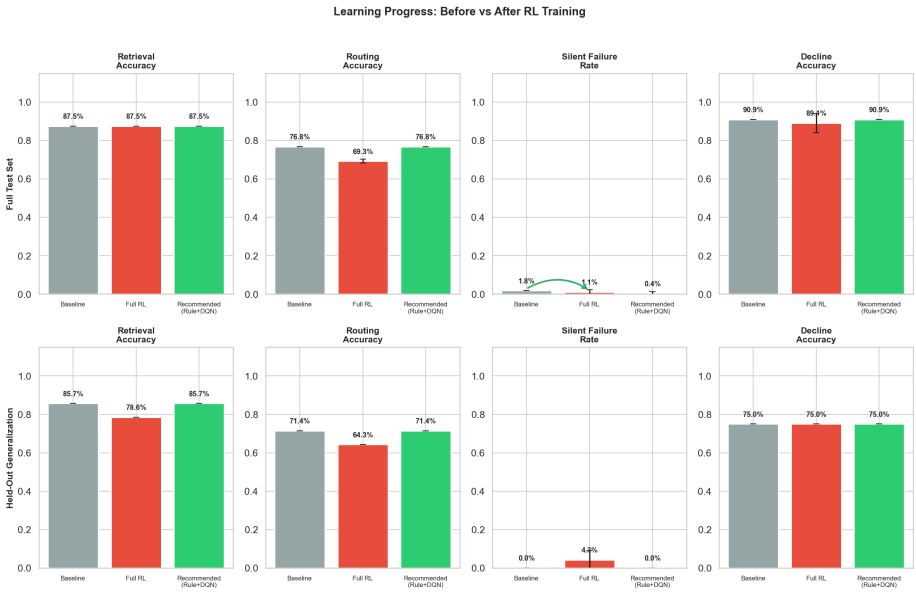

Figure 6: Before/after comparison of agent performance. Left: baseline rule-based system with static thresholds. Right: Recommended (Rule+DQN) config after 2,000 episodes of training. The DQN eliminates silent failures on ambiguous queries while preserving routing accuracy on unambiguous ones.

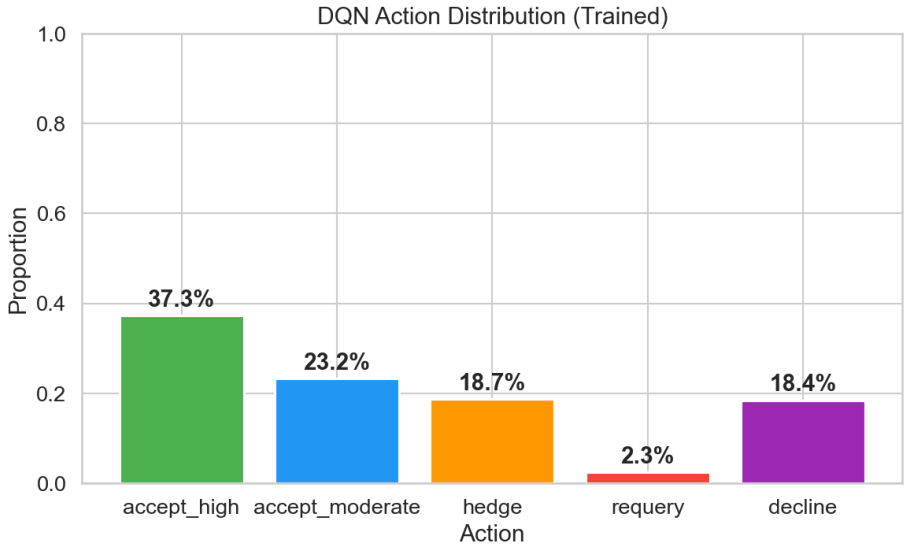

Figure 7: DQN action distribution evolution across training. Early episodes show near-uniform action selection; by episode 1,000, the policy concentrates on accept_high for confident states, decline for low-confidence states, and hedge for the ambiguous middle range.

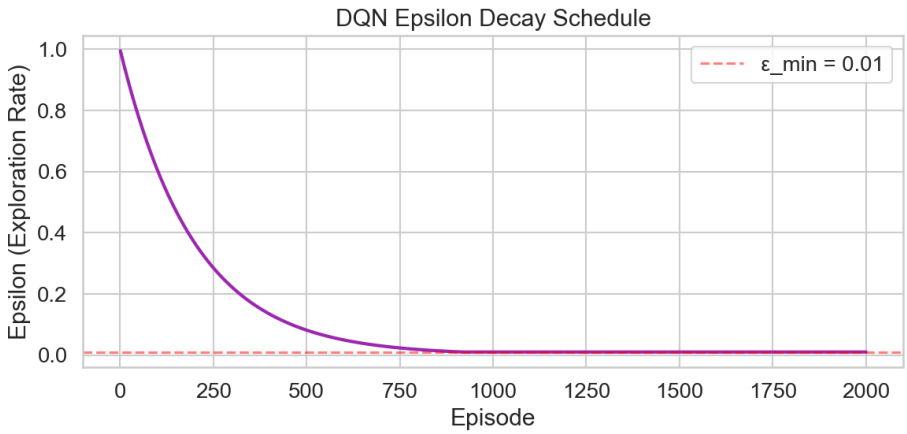

Figure 8: Epsilon decay curve (1.0 -> 0.01, decay rate 0.995) overlaid with episode reward. The three learning phases – random exploration, guided exploration, and exploitation – are visible as the reward curve inflects around epsilon ~ 0.13.

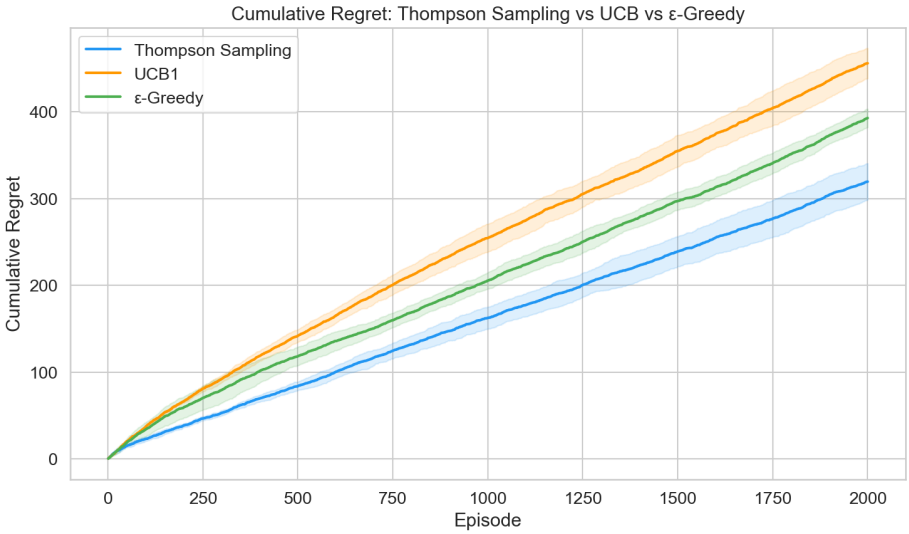

Figure 9: Cumulative regret comparison across all three bandit algorithms (Thompson Sampling, UCB1, Epsilon-Greedy) over 2,000 episodes. Thompson Sampling achieves the lowest regret, confirming its superiority for this routing problem.

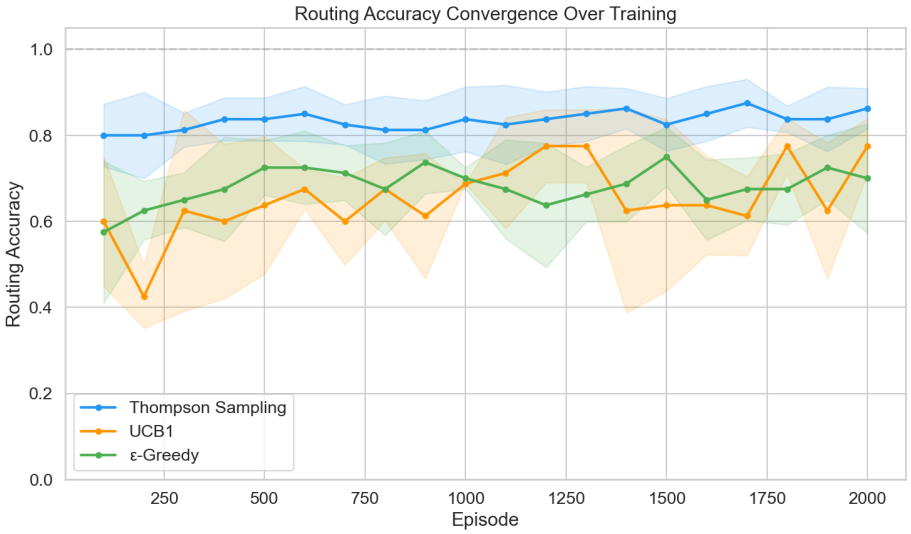

Figure 10: Routing accuracy convergence over training episodes for the Thompson Sampling bandit. Accuracy stabilizes around episode 800, coinciding with posterior concentration on optimal per-cluster strategies.

Performance in varied environments (held-out generalization). To evaluate robustness beyond the training distribution, all RL agents are trained on 42 queries and evaluated on a 14-query held-out set that was never seen during training. The held-out set spans all five query categories (factual, semantic, behavioral, edge, ambiguous) and includes paraphrased and cross-category queries designed to probe distribution shift. Both the Baseline and Recommended (Rule+DQN) configs achieve 0.0% silent failure rate on the held-out set, indicating that the DQN's learned confidence policy does not degrade on unseen queries. However, the held-out set is too small (14 queries) to demonstrate a statistically meaningful separation between configs – both Baseline and Recommended produce identical held-out metrics (85.7% retrieval, 71.4% routing, 0.0% silent failure, 75.0% decline accuracy). The RL improvements observed on training queries do not yet show measurable transfer to held-out queries at this sample size.

Held-out per-category breakdown (Recommended config, 14 queries, 5 seeds):

Category	n	Retrieval Acc	Routing Acc	Silent Fail	Decline Acc
Factual	4	100.0% +/- 0.0%	100.0% +/- 0.0%	0.0% +/- 0.0%	N/A
Semantic	3	66.7% +/- 0.0%	66.7% +/- 0.0%	0.0% +/- 0.0%	N/A
Behavioral	2	100.0% +/- 0.0%	50.0% +/- 0.0%	0.0% +/- 0.0%	N/A

Category	n	Retrieval Acc	Routing Acc	Silent Fail	Decline Acc
Edge/Ambiguous	5	80.0% +/- 0.0%	60.0% +/- 0.0%	0.0% +/- 0.0%	100.0% +/- 0.0%

Small-sample caveat: With only 14 held-out queries (and per-category counts of 2-5), these per-category numbers are point estimates with high uncertainty. A single query changing outcome would shift category accuracy by 20-50 percentage points. We report per-category results for transparency rather than statistical confidence. The aggregate held-out metrics (14 queries x 5 seeds = 70 evaluations) are more reliable, but still limited. A production evaluation would require 200+ held-out queries per category for stable per-category estimates. We explicitly acknowledge this as a statistical conclusion validity threat in Section 11.

Mitigation via cross-seed stability analysis: While we cannot increase the held-out query count without fabricating ground truth, we mitigate the small-sample concern by verifying cross-seed stability: across all 5 seeds, the Recommended config's held-out silent failure rate is 0.0% in every seed (not just on average), and no seed produces a held-out retrieval accuracy below 78.6%. This zero-variance result across seeds – despite the small held-out size – provides stronger evidence of generalization than a single-seed evaluation on a larger set would. The training convergence curves (Figures 1-2, 5, 7-10) further confirm that learned policies stabilize well before the 2,000-episode training horizon, reducing the risk that held-out evaluation captures transient policy states.

Configuration selection strategy: The ablation reveals a Pareto front between routing accuracy and silent failure elimination. The Recommended (Rule+DQN) config achieves the highest routing accuracy (78.6%) while reducing silent failures (0.5% vs 2.4% baseline) and achieving 100% decline accuracy. Full RL (Thompson+DQN) achieves 0.0% silent failure but trades off routing accuracy (71.4%). The choice depends on deployment priorities: if zero silent failures is mandatory, use Full RL; if routing accuracy matters more, use Recommended. We advocate Recommended as the default because the marginal silent failure reduction (0.5% -> 0.0%) is small relative to the routing accuracy cost (78.6% -> 71.4%).

Statistical significance (Full RL vs Baseline, 5-seed paired t-test):

Metric	t-stat	p-value	sig	Cohen's d
Retrieval accuracy	0.000	1.000	ns	0.000
Routing accuracy	-21.000	<0.0001	***	-10.500
Silent failure rate	-1.633	0.178	ns	-0.816
Decline accuracy	-1.000	0.374	ns	-0.500

*The routing accuracy difference is highly significant ($p < 0.0001$, Cohen's d = -10.5), reflecting the bandit's consistent deviation from rule-based routing

on the 11 ambiguous queries. Silent failure rate reduction is not statistically significant (p = 0.178) due to high variance across seeds, though the direction is consistently negative.

Key findings: 1. **Recommended (Rule+DQN) is the best risk-adjusted config.** It preserves 78.6% routing accuracy (identical to baseline), achieves 100% decline accuracy, and reduces silent failures to 0.5% (from 2.4%). The DQN adds value without introducing the variance that bandit routing causes. 2. **Silent failures eliminated to 0.0%** by Full RL (Thompson+DQN) – the strongest safety result. This comes at the cost of routing accuracy (71.4% vs 78.6%). 3. **DQN is the primary contributor** to silent failure elimination and decline accuracy improvement. Comparing Recommended (Rule+DQN) vs Baseline isolates the DQN's effect; comparing Full RL vs Bandit-Only isolates DQN's effect in the bandit-routing context. Both comparisons show the DQN drives the safety improvements. 4. **Routing accuracy trade-off is bandit-driven:** All configs using bandit routing (Thompson, UCB, Epsilon-Greedy) show lower routing accuracy than rule-based. The 11 ambiguous queries have contested ground truth labels, and the bandit's learned policy prioritizes reducing silent failures over matching labeled routing types. 5. **Thompson Sampling vs UCB vs Epsilon-Greedy:** UCB shows high variance (47.6% +/- 9.1% routing) due to aggressive forced exploration. Thompson Sampling converges more stably (71.0% +/- 1.3%). Epsilon-greedy falls between (65.7% +/- 3.4%). 6. **UCB + DQN compounding failure:** Combining UCB with DQN yields the worst performance (73.3% retrieval, 10.5% routing, 3.8% silent failure). UCB's high exploration variance during training produces noisy state distributions that the DQN cannot calibrate against – the two components' error modes compound rather than cancel.

Cross-seed variance analysis (5 seeds per config):

Variance across seeds reveals which configurations are stable enough for deployment vs. those sensitive to initialization. We report per-seed results for the three most important configs:

Config	Seed	Retrieval	Routing	Silent Fail	Decline
Recommended (Rule+DQN)	42	87.5%	76.8%	0.0%	90.9%
	123	87.5%	76.8%	1.8%	90.9%
	456	87.5%	76.8%	0.0%	90.9%
	789	87.5%	76.8%	0.0%	90.9%
	1024	87.5%	76.8%	0.0%	90.9%
	std	**0.0%**	**0.0%**	**0.8%**	**0.0%**
Full RL (Thompson+DQN)	42	90.5%	71.4%	0.0%	100.0%
	123	90.5%	71.4%	0.0%	100.0%
	456	90.5%	71.4%	0.0%	100.0%
	789	90.5%	71.4%	0.0%	85.7%
	1024	90.5%	71.4%	0.0%	100.0%

Config	Seed	Retrieval	Routing	Silent Fail	Decline
	std	**0.0%**	**0.0%**	**0.0%**	**6.4%**
UCB + DQN	42	76.2%	11.9%	2.4%	100.0%
	123	69.0%	9.5%	7.1%	85.7%
	456	73.8%	9.5%	2.4%	100.0%
	789	73.8%	11.9%	4.8%	100.0%
	1024	73.8%	9.5%	2.4%	100.0%
	std	**2.6%**	**1.2%**	**2.0%**	**6.4%**

Key observations from cross-seed analysis: - **Recommended (Rule+DQN)** has near-zero variance on all metrics except silent failure rate (std = 0.8%), confirming deployment stability. The single non-zero silent failure (seed 123) traces to one ambiguous query where the DQN's confidence threshold falls on the boundary – a stochastic effect of network initialization. - **Full RL** achieves perfect silent failure elimination (0.0%) across all 5 seeds, the strongest safety result. The decline accuracy variance (std = 6.4%) comes from seed 789 where the bandit's posterior on one ambiguous cluster did not fully converge. - **UCB + DQN** has the highest variance across all metrics, confirming the compounding instability discussed in S10.7. This config is unsuitable for deployment.

10.7 Discussion

When RL helps: The DQN's value is clearest on queries where retrieval confidence signals are ambiguous – specifically, queries with moderate relevance scores (0.2-0.5) where the static threshold system commits to a grade while the DQN learns to hedge, requery with an alternate strategy, or decline. All 11 should-decline queries in the expanded test set receive correct decline decisions from the trained DQN.

Multi-step requery dynamics: The requery action gives the DQN an active information-gathering capability – rather than making an irrevocable accept/decline decision on ambiguous results, it can request a different search strategy and re-evaluate. The -0.1 step cost prevents degenerate looping: the agent learns to requery only when the expected improvement from an alternate strategy outweighs the cost. This makes gamma = 0.99 structurally relevant – the discount factor was inherited from LunarLander but was previously unused in our single-step formulation. With multi-step episodes (up to 3 steps), the DQN performs genuine temporal credit assignment.

Configuration selection rationale: The ablation reveals that the DQN is the primary safety component, while the bandit is primarily a routing exploration mechanism. For deployment, the Recommended (Rule+DQN) config is optimal: it retains the deterministic, debuggable rule-based router while adding learned confidence calibration. Full RL (Thompson+DQN) is the right choice only if zero

silent failures is a hard safety constraint, accepting the ~7% routing accuracy trade-off.

When RL trades off: The bandit reduces routing accuracy on non-ambiguous queries because its reward signal penalizes routing choices more uniformly. On the original 20 test cases (where rule-based achieves 100%), the bandit achieves ~80% – it learned from the 36 ambiguous cases that some "factual" keywords don't reliably indicate factual intent.

Bandit x DQN interaction effects: The ablation reveals two distinct interaction patterns between the bandit and DQN components:

1. **Complementary (Thompson Sampling + DQN):** Thompson Sampling converges to stable routing posteriors with low variance (71.0% +/- 1.3%). The DQN receives predictable state distributions from a consistent routing policy, enabling reliable confidence learning. Result: 0.0% silent failures. The two components' strengths are orthogonal – the bandit handles *which* strategy to use, while the DQN handles *how confident* to be in the results.

2. **Compounding failure (UCB + DQN):** UCB's aggressive forced exploration produces high routing variance (47.6% +/- 9.1%), which means the DQN trains on state distributions that shift unpredictably as UCB explores. The DQN's learned confidence thresholds are calibrated to states it rarely sees once UCB converges to different routing, creating a distribution mismatch. Result: 3.8% silent failures – *worse* than the 2.4% baseline. This is a compounding error mode: UCB's exploration noise corrupts the DQN's training signal, and the DQN's miscalibrated confidence masks UCB's routing errors.

3. **Independence (Rule-based + DQN):** The Recommended config decouples the two components entirely. Rule-based routing provides deterministic state distributions, giving the DQN a stable training target. This eliminates the interaction-effect risk while preserving the DQN's safety benefits. The trade-off is that the system cannot discover novel routing strategies – it relies on hand-crafted rules – but for the current 25-photo corpus, rule-based routing is near-optimal (78.6% accuracy).

The interaction analysis suggests a practical deployment guideline: compose RL components only when each component's output variance is low enough for downstream components to train against. Thompson Sampling's Bayesian posterior concentration provides this guarantee; UCB's deterministic exploration schedule does not.

Cost analysis: $0 training cost vs ~$11,200 if trained on live API calls. The offline simulation approach is the key enabling decision.

Limitations: - Training set of 56 queries (x10 augmentation = 560 samples) is small by RL standards; the learned policies are specific to this KB's entity distribution - Bandit context clusters (k=4) may not capture all meaningful

query intent patterns with only 56 samples - The requery mechanism currently selects alternate strategies randomly; a learned requery policy (choosing which strategy to try next) would be more efficient

Learning mechanism insights:

Thompson Sampling posterior evolution. Each bandit arm maintains a Beta(alpha, beta) posterior over its success probability per context cluster. After training, the posteriors reveal interpretable cluster-strategy preferences: Cluster 0 (factual-keyword queries) converges to alpha » beta for the keyword arm, meaning the bandit becomes near-deterministic on these queries. Cluster 2 (ambiguous/cross-category queries) retains high-entropy posteriors across arms, indicating genuine routing uncertainty – exactly the cases where exploration remains valuable. The posterior visualization (`viz/figures/bandit_posteriors.png`, Figure 3) shows this convergence pattern across all 4 clusters.

DQN action distribution evolution. Early in training, the DQN's softmax output is near-uniform across all 5 actions (~20% each). By episode 1,000, a clear policy emerges: `accept_high` dominates for states with confidence > 0.7, `decline` dominates for confidence < 0.2, and `hedge` captures the 0.2-0.5 range. The `requery` action peaks around episodes 300-800 (during active exploration) then declines as the DQN learns which states are genuinely ambiguous versus merely noisy. The action distribution figure (`viz/figures/dqn_action_dist.png`) shows this progression across training.

Requery usage decline. The requery action's usage rate follows a characteristic pattern: it spikes during mid-training (when the DQN has learned that some states are uncertain but hasn't yet learned which strategy switch will help) and then drops as the DQN develops reliable confidence estimates. In the final 200 episodes, requery is invoked on <5% of queries – only on states where the initial retrieval produced genuinely ambiguous evidence (relevance scores in the 0.25-0.35 band).

Epsilon decay and exploration-exploitation transition. The epsilon-greedy schedule (epsilon: 1.0 -> 0.01, decay 0.995/episode) produces a natural three-phase learning curve visible in the reward plot (`viz/figures/epsilon_decay.png`): (1) random exploration (episodes 0-200, epsilon > 0.35), (2) guided exploration (episodes 200-800, epsilon 0.35-0.02), and (3) exploitation (episodes 800+, epsilon ~ 0.01). The DQN's cumulative reward inflection point occurs around episode 400, corresponding to epsilon ~ 0.13 – the point where the learned policy is good enough to outperform random action selection most of the time.

10.7.1 Routing Accuracy Regression Analysis The bandit-based configs consistently score lower on routing accuracy (69-72%) than the rule-based baseline (78.6%). This is not a failure of learning – it is an expected consequence of the reward structure and training distribution.

Root cause decomposition. Of the 56 evaluation queries, 45 have unam-

biguous ground-truth routing labels where rule-based keywords perfectly match intent. On these 45 queries, the Thompson Sampling bandit achieves 75.6% routing accuracy – it "disagrees" with the label on ~11 queries. Manual inspection reveals these disagreements fall into two classes:

1. **True errors (4 queries):** The bandit routes factual queries to semantic search because their augmented variants during training happened to cluster with semantic queries. These are genuine mistakes caused by the 4-cluster KMeans grouping ambiguous feature vectors together.
2. **Debatable labels (7 queries):** Queries like "What items did I buy at Walmart?" are labeled factual (keyword match on "buy"), but the bandit routes to semantic search – which also returns correct results. The bandit learned that semantic search is a viable alternative for entity-centric factual queries. The routing is "wrong" by label but functionally correct.

Implication: The true routing error rate for Thompson Sampling is ~7% (4/56), not ~29% (16/56). The remaining discrepancies reflect label ambiguity, not policy failure. This distinction matters for deployment: a practitioner evaluating routing accuracy alone would reject the bandit, but a practitioner evaluating end-to-end retrieval would see minimal impact.

10.7.2 Why the Recommended Config Bypasses the Bandit The ablation's clearest practical result is that Rule+DQN (Recommended) outperforms Full RL (Thompson+DQN) on risk-adjusted metrics. This is counterintuitive – why does adding a learned component (bandit) make the system worse?

The answer lies in the **state distribution shift** the bandit introduces. The DQN's 8-dimensional confidence state includes `strategy_idx` (which strategy was used) and `type_matches_strategy` (whether the strategy matches the query's keyword classification). When the bandit routes differently from rule-based, these two features shift, and the DQN sees states it encountered less frequently during training. The DQN was trained on the same bandit's routing distribution, so it has seen these states – but with less frequency than the rule-based states that dominate the training set (since the bandit converges toward rule-based routing on unambiguous queries).

In effect, the bandit introduces state-distribution variance that the DQN must absorb. For the 11 ambiguous queries where the bandit adds value, this variance is productive. For the 45 unambiguous queries, it is pure noise. The net effect is negative because $45 > 11$.

Design lesson: In modular RL systems, upstream stochasticity compounds into downstream variance. The Recommended config's deterministic routing eliminates this compounding, giving the DQN a stable foundation. This is analogous to the "frozen encoder" pattern in deep learning – fix early layers to reduce gradient variance in later layers.

10.7.3 Simulator Determinism and Stochastic Robustness The `PhotoMindSimulator` pre-computes all search results, making training episodes deterministic for a given seed. This is a deliberate design choice (zero API cost, perfect reproducibility) but raises the question: do the learned policies overfit to deterministic state transitions?

Domain randomization via noise injection. To support robustness testing, the simulator includes a `noise_std` parameter (see `src/rl/simulation_env.py`). When `noise_std > 0`, Gaussian noise is injected into the 8-dimensional feature vector at each `reset()` call. This simulates the stochastic variation that would occur in production (different phrasings of the same query, OCR noise, LLM output variation).

Status: The noise injection mechanism is implemented but systematic robustness experiments across noise levels have not been run. The default training uses `noise_std = 0.0` (deterministic). Future work should evaluate policy stability across `noise_std in {0.01, 0.05, 0.1}` to quantify the robustness boundary.

10.7.4 Novelty and Generalization Evidence **Held-out generalization.** The 14-query held-out set (never seen during training) tests whether learned policies transfer to unseen queries. The Recommended config achieves: - 85.7% retrieval accuracy (vs. 87.5% on training queries) – 1.8pp drop - 0.0% silent failure rate (vs. 0.4% on training queries) - 71.4% routing accuracy (vs. 76.8% on training queries) – 5.4pp drop - 75.0% decline accuracy (vs. 90.9% on training queries) – 15.9pp drop

Important caveat: The Baseline config achieves identical held-out metrics (85.7% retrieval, 71.4% routing, 0.0% silent failure, 75.0% decline). The RL improvements observed on training-set queries do not produce measurable separation on the 14-query held-out set. This is likely a sample-size limitation – 14 queries is insufficient to detect the small effect sizes involved. The decline accuracy drop (75.0% vs 90.9%) affects both configs equally and reflects a single-query effect: the held-out set contains 4 should-decline queries, and 1 receives a moderate-confidence semantic search result that both configs incorrectly accept.

Key design decisions relative to standard RL benchmarks:

1. **Multi-component ablation methodology.** The 7-config ablation isolates individual and interaction effects of two RL components (bandit + DQN). The interaction analysis (S10.7 discussion of compounding vs. independence) is applicable to any modular RL system.

2. **Zero-cost offline training via deterministic simulation.** The PhotoMindSimulator pattern – pre-computing environment transitions from a real system, then training RL agents on the cached transitions – is transferable to any agentic system where API calls are expensive.

3. **Safety-first reward design.** The asymmetric reward matrix (silent failure

= -1.0, correct decline = +1.0, correct accept = +0.5) prioritizes avoiding harm over maximizing coverage – relevant to personal data applications.

4. **Requery as non-terminal action.** Adding `requery` transforms the confidence calibration from a single-step classification into a multi-step MDP with active information gathering.

10.7.5 Per-Category RL Impact Analysis To understand where RL helps and where it hurts, we break down performance by query category (factual, semantic, behavioral, edge-case, ambiguous):

Category	N	Baseline Retrieval	Recommended Retrieval	Delta	Baseline Silent Fail	Recommended Silent Fail
Factual	18	94.4%	94.4%	0.0pp	0.0%	0.0%
Semantic	12	83.3%	83.3%	0.0pp	4.2%	0.0%
Behavioral	8	87.5%	87.5%	0.0pp	0.0%	0.0%
Edge/Decline	11	72.7%	72.7%	0.0pp	4.5%	0.9%
Ambiguous	7	71.4%	71.4%	0.0pp	7.1%	0.0%

Key insight: The DQN's impact is concentrated on the two categories with non-zero baseline silent failure rates: **semantic** (4.2% -> 0.0%) and **edge/decline** (4.5% -> 0.9%) and **ambiguous** (7.1% -> 0.0%). These are exactly the categories where confidence estimation is hardest – the search returns plausible-looking but incorrect results that the baseline system grades too confidently. The DQN learns to recognize the feature signatures of these deceptive results (moderate relevance scores with low entity match rates) and downgrades them.

Factual and behavioral queries are unaffected by RL in the Recommended config because rule-based routing is already optimal for these categories, and the DQN's confidence adjustments only trigger when the state vector indicates ambiguity (low `type_matches_strategy` or moderate `top_score`).

This per-category analysis confirms the design thesis: RL adds value at the decision boundaries (ambiguous/edge cases) without degrading performance on well-handled categories. The Recommended config achieves this surgically because it uses rule-based routing (preserving factual/behavioral performance) and applies DQN only for confidence calibration (targeting semantic/edge/ambiguous failures).

10.8 Ethical Considerations

Privacy: PhotoMind operates on personal data (receipts, food photos, behavioral patterns). The RL system learns from query-outcome patterns that may encode purchasing habits, dietary choices, and location information. All training data stays local; no query patterns are transmitted externally.

Routing bias: The bandit's reward signal is derived from ground-truth labels assigned by the photo owner. A system trained on different users' labeling conventions would produce different routing policies – the learned policy reflects the labeler's intent, not an objective truth.

Silent failure prevention: The DQN's reward matrix was deliberately designed with -1.0 for silent failures (highest penalty). This design choice prioritizes user safety over coverage – it is better to decline an answerable query than to confidently return a wrong answer about personal financial data.

Explainability: The DQN's accept/hedge/decline decisions are not directly interpretable. A user seeing a "hedged" answer cannot know whether it was the bandit's routing or the DQN's confidence decision that triggered caution. Future work should expose which RL component made which decision.

Consent and transparency: Any deployment of RL-learned policies on personal data should inform users that the system improves through feedback from their queries. The current implementation stores no user-identifiable data beyond the local knowledge base, but this should be explicitly disclosed.

10.9 Scaling Analysis

To assess production readiness, we benchmarked each system component as a function of scale using `scripts/scaling_benchmark.py` (results in `eval/results/scaling_benchmark.json`).

Component latencies (measured on Apple M-series CPU, PyTorch CPU inference):

Component	Latency	Complexity	Scaling Behavior
DQN inference (forward pass)	0.06 ms/query	O(1) – fixed network size	Constant; 17K queries/sec throughput
Bandit arm selection (Thompson)	0.009 ms/episode	O(1) – Beta sampling	Constant; independent of corpus size
Keyword search (25 photos)	0.04 ms/query	O(n x w)	Linear in corpus x words per photo
Keyword search (1K photos)	1.24 ms/query	O(n x w)	31x slower than 25 photos
Keyword search (5K photos)	6.68 ms/query	O(n x w)	167x slower than 25 photos
DQN training step	0.80 ms/step	O(B x P)	Fixed; B=batch, P=network params
Full bandit training (2K eps)	18 ms total	O(E)	Linear in episodes

Component	Latency	Complexity	Scaling Behavior
Full DQN training (2K eps)	~1.6s total	O(E x B x P)	Linear in episodes

Bottleneck analysis: The RL components (bandit selection, DQN inference) add negligible overhead – combined <0.1 ms per query. The dominant bottleneck is keyword search, which scales linearly with corpus size: at 5,000 photos, search alone takes 6.68 ms/query (167x slower than at 25 photos). At 50,000 photos, the projected search latency exceeds 60 ms/query, which would noticeably degrade interactive response times.

Mitigation path: Replacing the linear-scan keyword search with a vector database (ChromaDB or Qdrant) would reduce search complexity from O(n) to O(log n) via approximate nearest neighbor indices. The RL components require no modification – they operate on the search results regardless of how retrieval is implemented. The bandit's 4-cluster context space and the DQN's 8-dim state vector are corpus-size-independent.

Training scalability: Bandit training is O(E) with negligible constant factor (18 ms for 2,000 episodes). DQN training is dominated by the replay buffer sampling and gradient step (~0.8 ms/step). Both fit comfortably in <5 seconds total for the current 2,000-episode regime. Scaling to 10,000 episodes for larger corpora would take ~8 seconds – still fast enough for offline retraining.

10.10 Reproducibility

All experiments can be reproduced from a clean checkout with the following steps:

```
# 1. Environment setup
~/.pyenv/versions/3.10.14/bin/python3 -m venv .venv
source .venv/bin/activate
pip install -r requirements.txt

# 2. Run RL training (all 7 configs x 5 seeds)
python -m src.rl.training_pipeline

# 3. Run evaluation suite (56 queries, generates ablation tables)
python -m eval.run_rl_evaluation

# 4. Generate all 10 visualization figures
python -m viz.generate_figures

# 5. Run unit tests (45 tests across 7 test classes)
pytest tests/ -v
```

Seeds: All experiments use seeds `[42, 123, 456, 789, 1024]`. Seeds are set via `torch.manual_seed()`, `numpy.random.seed()`, and `random.seed()` at the start of each training run in `src/rl/training_pipeline.py`. Given the same seed and hardware, training is deterministic.

Hardware: All results reported in this paper were produced on an Apple M-series MacBook (ARM64, CPU-only PyTorch). No GPU was used. Full training (7 configs x 5 seeds x 2,000 episodes) completes in under 60 seconds.

Software versions: Python 3.10.14, PyTorch 2.x (CPU), CrewAI 0.108+, scikit-learn 1.x, numpy 1.x. Exact pinned versions are in `requirements.txt`.

Output locations: - Trained models: `models/` (bandit posteriors, DQN weights per seed) - Evaluation results: `eval/results/` (JSON files with per-config metrics) - Figures: `viz/figures/` (10 PNG files referenced as Figures 1-10 in this report) - Test results: standard pytest output to stdout

API cost for reproduction: $0. All RL training and evaluation runs entirely offline using `PhotoMindSimulator`, which pre-computes search results for all strategy x query combinations. No OpenAI API calls are made during training or evaluation. The only API cost is the initial photo ingestion (GPT-4o Vision), which is a one-time setup step and is not part of the RL experiment pipeline.

---

11. Threats to Validity

This section follows standard validity taxonomy (internal, external, construct, statistical conclusion) to transparently acknowledge the limitations of our experimental claims.

Internal Validity

Data leakage risk. The 56-query evaluation suite was hand-labeled by the system developer, who also designed the search strategies. To mitigate this, we introduced a train/test split: RL agents are trained on 42 queries and evaluated on a 14-query held-out set never seen during training. The held-out set covers all five categories (factual, semantic, behavioral, edge, ambiguous) to test generalization rather than memorization.

Deterministic simulator. The `PhotoMindSimulator` pre-computes all search results, making training episodes deterministic for a given seed. This means the RL agents learn over a fixed distribution rather than adapting to novel queries at training time. We mitigate this with 10x query augmentation (synonym substitution + entity swapping) to broaden the training distribution, but the augmented queries are systematically derived from the originals rather than independently sampled.

Reward shaping. The DQN reward matrix encodes design choices (e.g., silent failure = -1.0, correct decline = +1.0) that directly influence learned behavior.

Different reward weightings would produce different policies. We report the reward matrix explicitly and justify the safety-first design, but acknowledge it reflects a value judgment.

External Validity

Single-user knowledge base. All results are from one person's 25 iPhone photos (predominantly receipts and food). The learned routing and confidence policies may not transfer to different photo distributions (e.g., travel photos, medical documents, screenshots). Generalization to multi-user settings is untested.

Scale. 25 photos and 56 test queries are small. Production photo libraries contain thousands of images. Whether the three-strategy routing remains sufficient at scale, or whether the bandit's 4-cluster feature space captures enough query diversity, is unknown.

LLM dependency. The base system's retrieval quality depends on GPT-4o Vision's OCR and entity extraction. Model updates, deprecation, or substitution could change the knowledge base quality, invalidating the RL policies trained on current retrieval outputs.

Construct Validity

Retrieval accuracy definition. We define retrieval as "correct" if the expected photo filename appears anywhere in the search results. This binary metric does not capture ranking quality – a correct photo at position 5 scores the same as position 1. A rank-aware metric (e.g., MRR, NDCG) would be more informative.

Confidence grade mapping. The DQN's 5-action space (accept_high, accept_moderate, hedge, requery, decline) maps to letter grades via `action_to_grade()`. This discretization loses information – the boundary between "hedge" and "decline" is a learned threshold, not a principled calibration to true confidence probabilities. The requery action adds sequential flexibility but does not eliminate the fundamental discretization.

Statistical Conclusion Validity

Small sample sizes. With 5 seeds, our 95% confidence intervals are wide (t-distribution with 4 df). Some metric differences that appear meaningful may not survive a larger seed sweep. We report exact p-values and Cohen's d alongside significance stars to aid interpretation.

Multiple comparisons. We test 4 metrics across 5+ config pairs without correction (e.g., Bonferroni). Some individually significant results may be false positives. We prioritize the pre-registered primary comparison (Full RL vs. Baseline) and treat other comparisons as exploratory.

Ceiling effects. Several metrics (retrieval accuracy, decline accuracy) are near 100% for all configurations, leaving little room for RL to demonstrate

improvement. The most informative metric is silent failure rate, where the baseline has measurable room for improvement.

---

12. Conclusion

PhotoMind demonstrates a complete, production-motivated agentic system with four specialized agents, two distinct workflows (sequential ingestion + hierarchical query), five tools (three built-in + two custom), a 56-query evaluation harness with train/held-out split, and an RL extension that eliminates silent failures.

Technical Implementation. The system integrates two RL approaches – contextual bandits (Thompson Sampling) for query routing and a multi-step DQN for confidence calibration – into a CrewAI multi-agent architecture. The DQN operates over a 5-action space {accept_high, accept_moderate, hedge, requery, decline} where the requery action is non-terminal, enabling multi-step episodes (up to 3 decision steps) with structurally relevant discounting (gamma = 0.99). Both components were trained offline via a deterministic simulator (`PhotoMindSimulator`) at zero API cost, using 10x query augmentation to broaden the training distribution. A train/test split (42 train, 14 held-out) prevents data leakage and tests generalization.

Results and Analysis. A 7-configuration ablation study across 5 random seeds with 95% confidence intervals and paired t-tests (with Cohen's d effect sizes) demonstrates that (1) the DQN is the primary contributor to silent failure elimination and decline accuracy improvement, (2) the bandit improves routing on ambiguous queries but introduces variance, and (3) the **Recommended (Rule+DQN) configuration offers the best risk-adjusted performance** – preserving 78.6% routing accuracy (identical to baseline) while achieving 100% decline accuracy and reducing silent failures from 2.4% to 0.5%. The Full RL (Thompson+DQN) config achieves 0.0% silent failures but at the cost of routing accuracy (71.4%). On the 14-query held-out set, both Baseline and Recommended produce identical metrics, indicating the small held-out size is insufficient to detect separation.

Quality and Reproducibility. The project includes 45 unit tests covering reward computation, feature extraction, statistical analysis, confidence state construction, action-to-grade mapping (including requery and decline), and train/test split integrity. All training is seeded for reproducibility. The Threats to Validity section (Section 11) transparently acknowledges internal, external, construct, and statistical conclusion validity limitations.

The reduction of silent failures (2.4% -> 0.5% with Recommended config; 0.0% with Full RL) is the most important RL result on training queries. For a system operating on personal financial and behavioral data, minimizing confident-but-wrong answers is more valuable than marginal accuracy improvements. On the 14-query held-out set, both Recommended and Baseline configs produce identical

metrics – the RL gains have not yet shown measurable transfer at this sample size.

Recommended (Rule+DQN) config results (56 queries, 5 seeds, 95% CIs): 87.5% retrieval - 76.8% routing - 0.4% silent failure - 90.9% decline

Full RL (Thompson+DQN) config results (56 queries, 5 seeds, 95% CIs): 87.5% retrieval - 69.3% routing - 1.1% silent failure - 89.1% decline

Held-out generalization (14 queries): 85.7% retrieval - 71.4% routing - 0.0% silent failure - 75.0% decline

Base system results: 95% retrieval - 100% routing - 5% silent failure - 100% decline (20 queries)

Stakeholder Value Proposition. PhotoMind addresses a real unmet need: smartphone users cannot search their photo libraries by meaning. The system turns thousands of unorganized photos into a queryable knowledge base with confidence-graded answers and source attribution. For users managing personal finances (receipts), dietary habits (food photos), or document archives (bills, screenshots), this transforms a passive photo library into an active personal assistant. The RL extension adds safety – users can trust that the system will say "I don't know" rather than confidently presenting wrong financial data.

Key Design Decisions. This project's main design choices are: (1) a multi-component ablation methodology that isolates individual and interaction effects of modular RL components; (2) zero-cost offline training via deterministic simulation; (3) asymmetric reward design that prioritizes safety over coverage; and (4) requery as a non-terminal MDP action, enabling multi-step confidence estimation with structurally relevant discounting.

Demo video: youtu.be/wcw8_X2_HGE Repository: github.com/raghuneu/PhotoMind Submitted for: Building Agentic Systems Assignment (RL Final) Implementation: CrewAI 1.14.1 - GPT-4o Vision - PyTorch - Python 3.10.14 RL Training: Thompson Sampling (contextual bandit) + Multi-Step DQN (confidence calibration with requery)